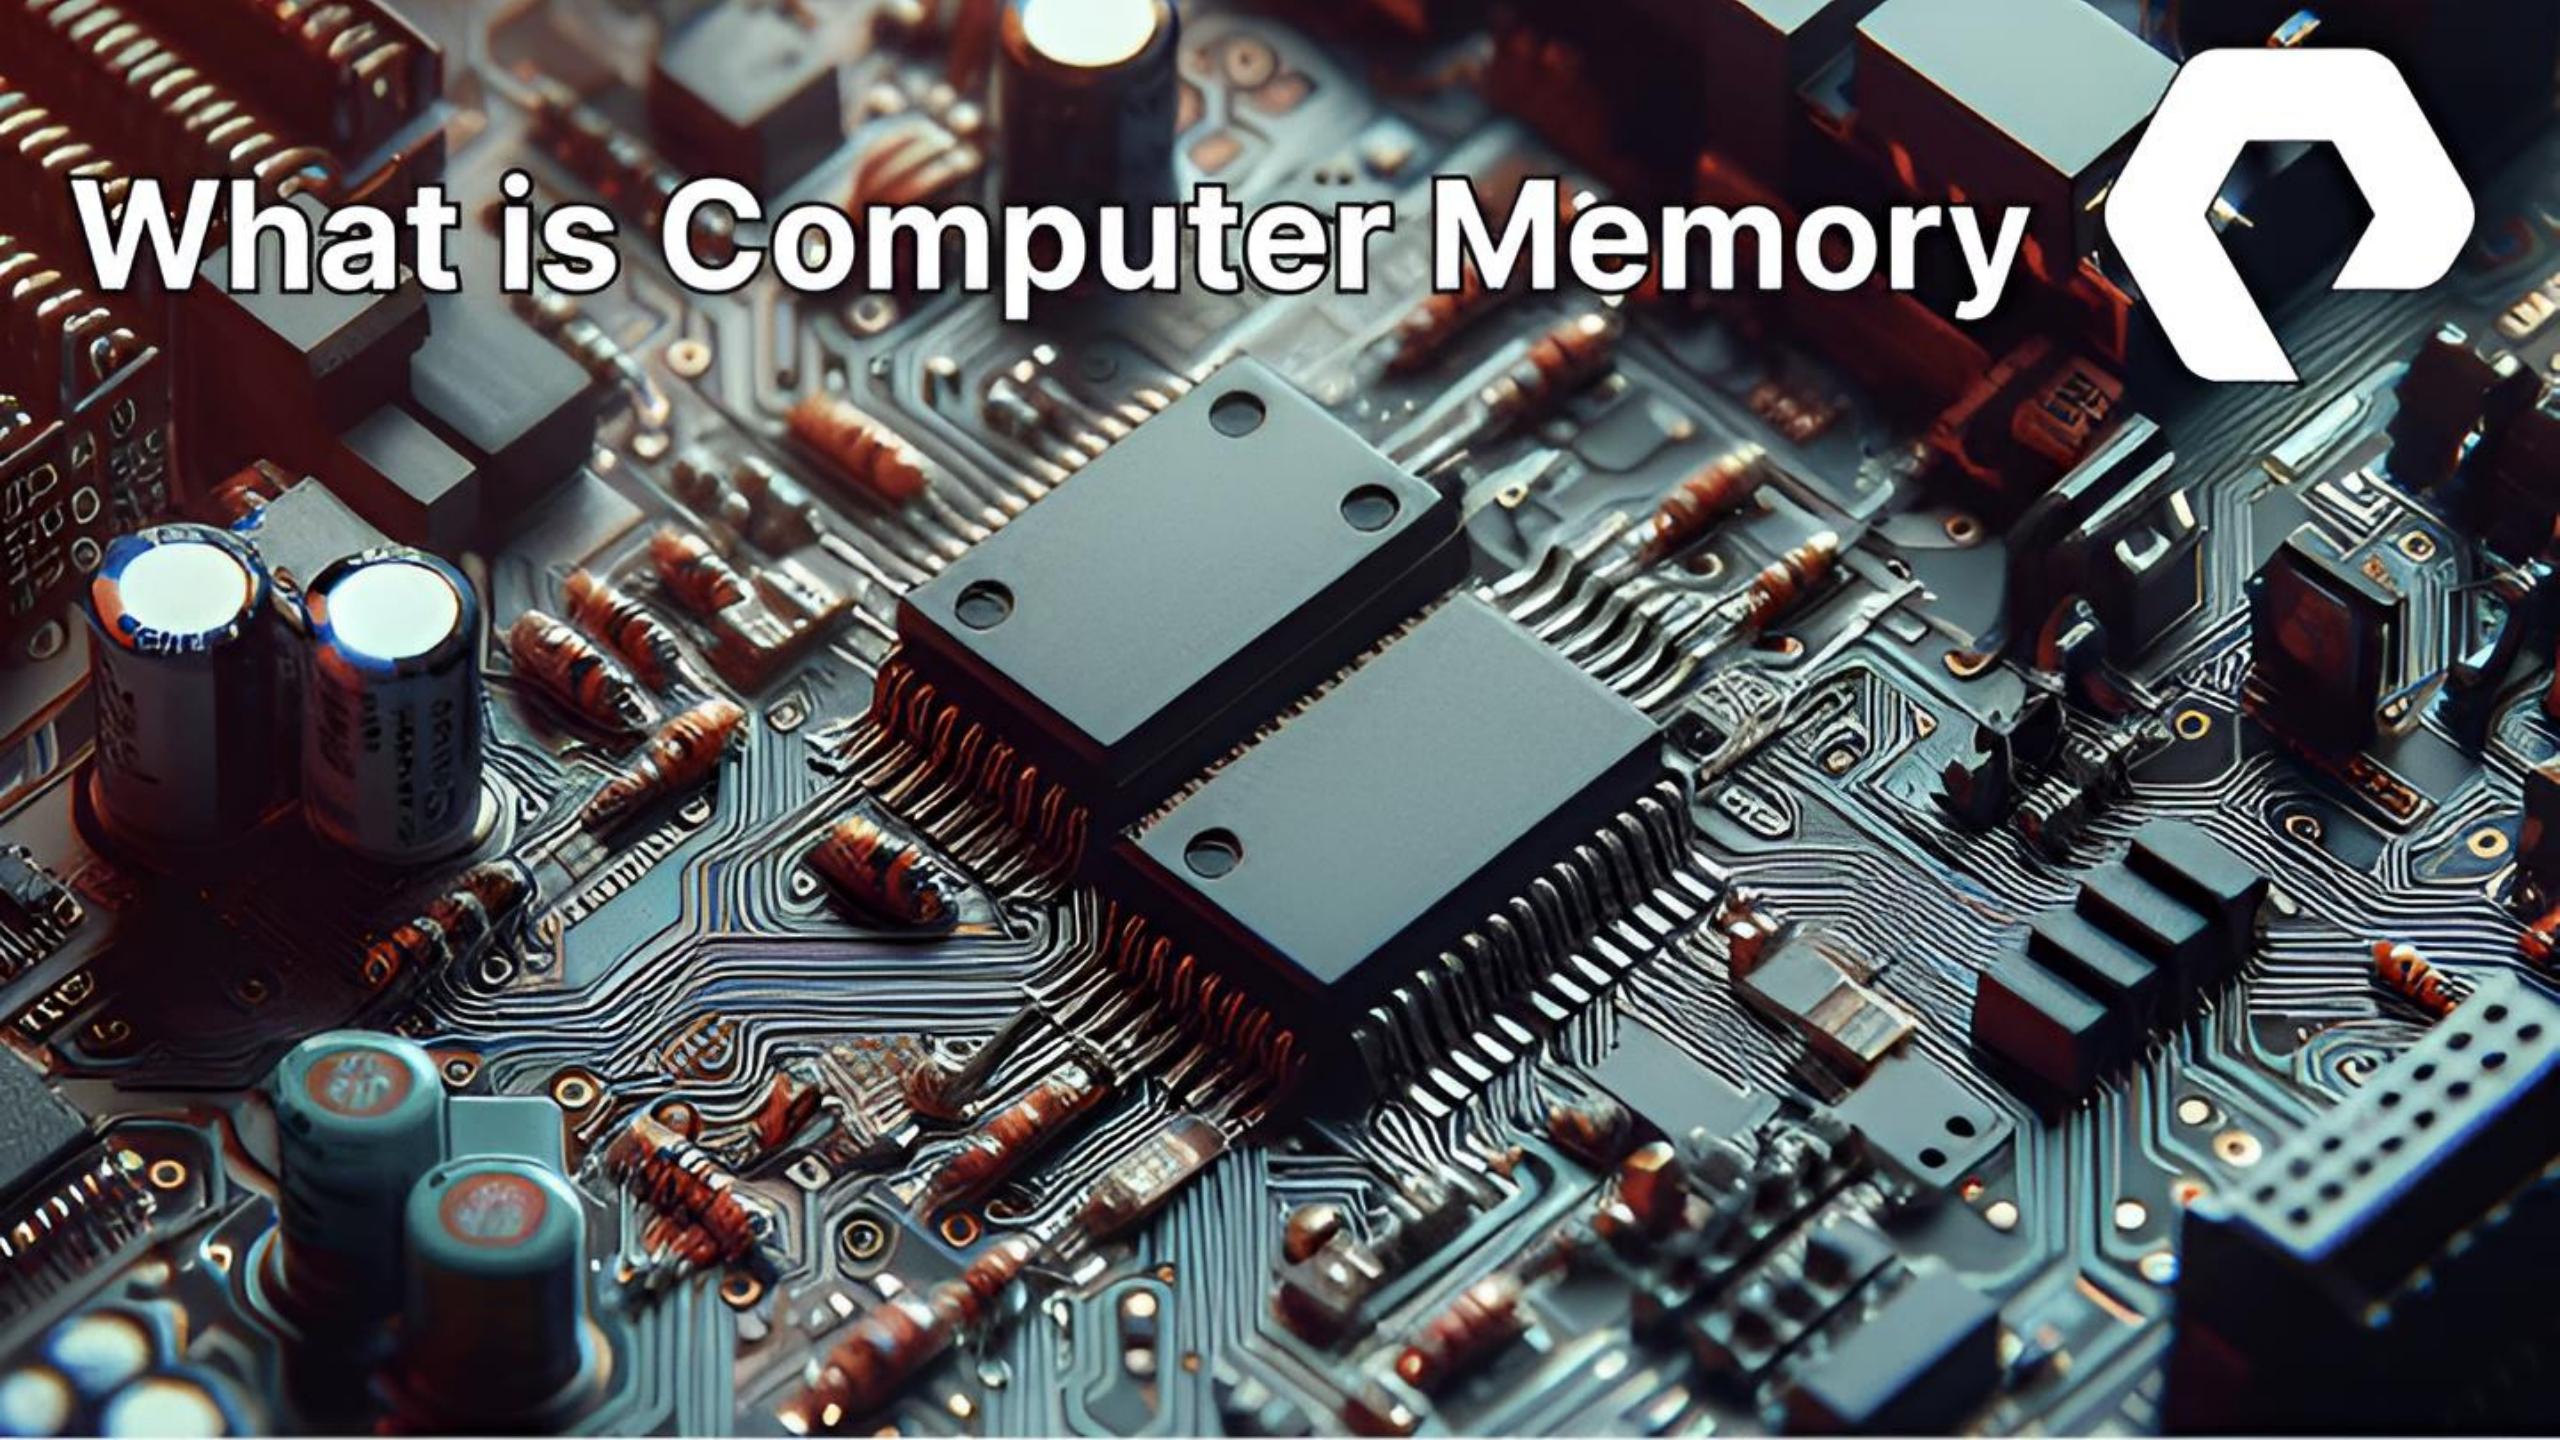
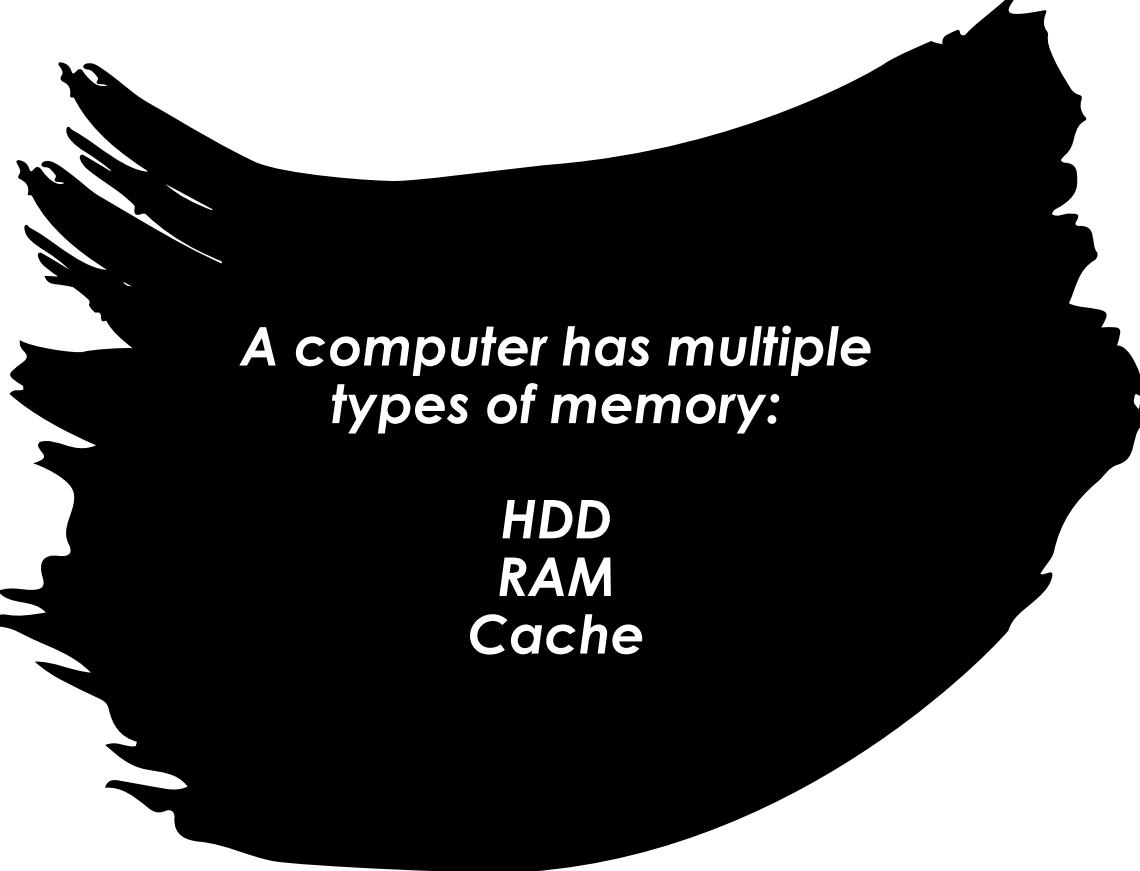


Optimization

- COMPUTERS MEMORY
- FAST CODE BUT HOW?
- OOP, IS IT ANY GOOD?
- STACK VS HEAP
- MEMORY POOL
- ALGORITHM & DATA STRUCTURE
- LEVEL OF DETAILS (LOD)
- NORMAL & PARALLAX MAPPING
- RENDER TO TEXTURES
- ATLAS TEXTURE VS SPRITE SHEET
- BATCHING
- LIGHT MAPS & PROBES
- OCCLUSION CULLING
- PREFABS
- PROFILING
- COMMON SENSE

What is Computer Memory

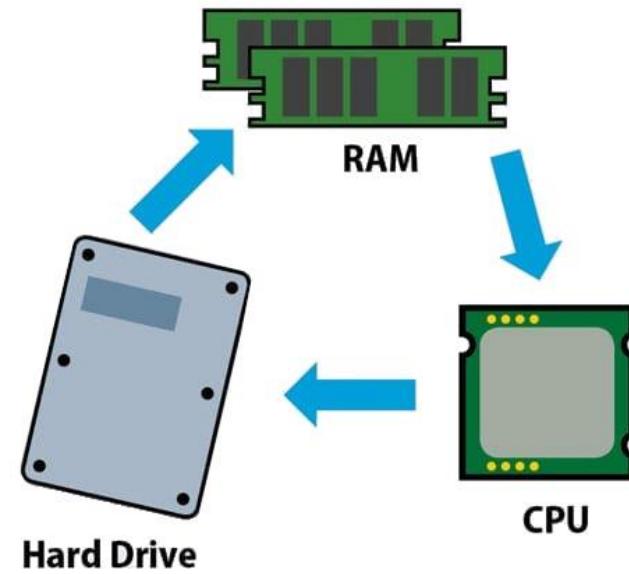




A computer has multiple types of memory:

**HDD
RAM
Cache**

Memory vs. Storage

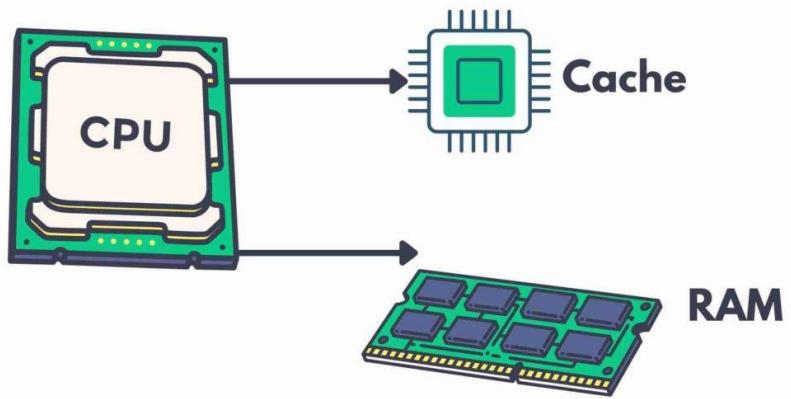


Memory

- Volatile
- Fast Access

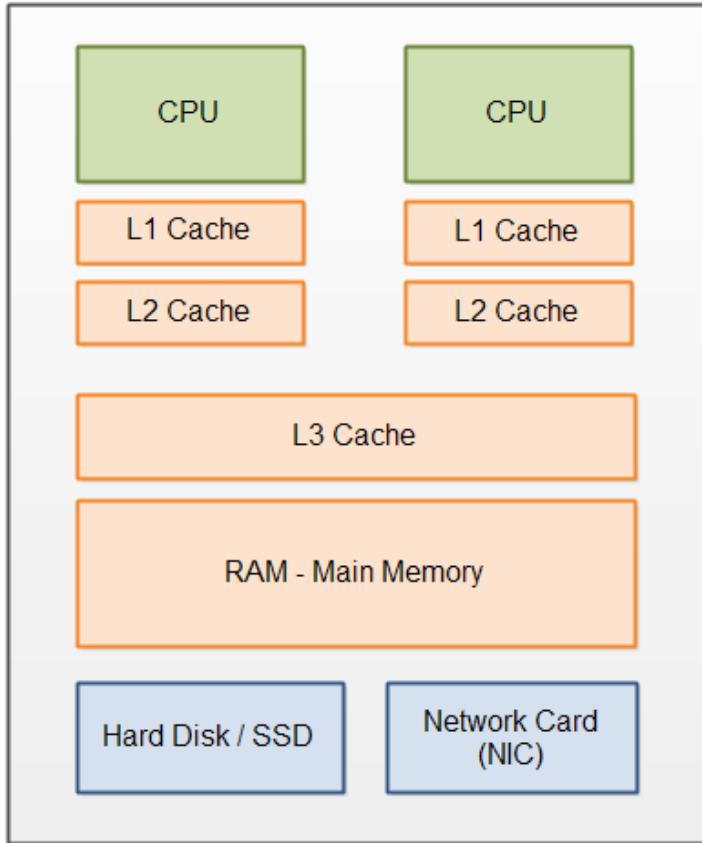
Storage

- Non-volatile
- Large Capacity
- Slower Access



What is memory cache?

- Bridge between RAM and CPU?
- Performance booster?
- Legacy from the past?



Memory Caches:

L3 - Fast

L2 - Faster

L1 - Fastest...

*smallest ...
closest to CPU*

Hierachal memory!

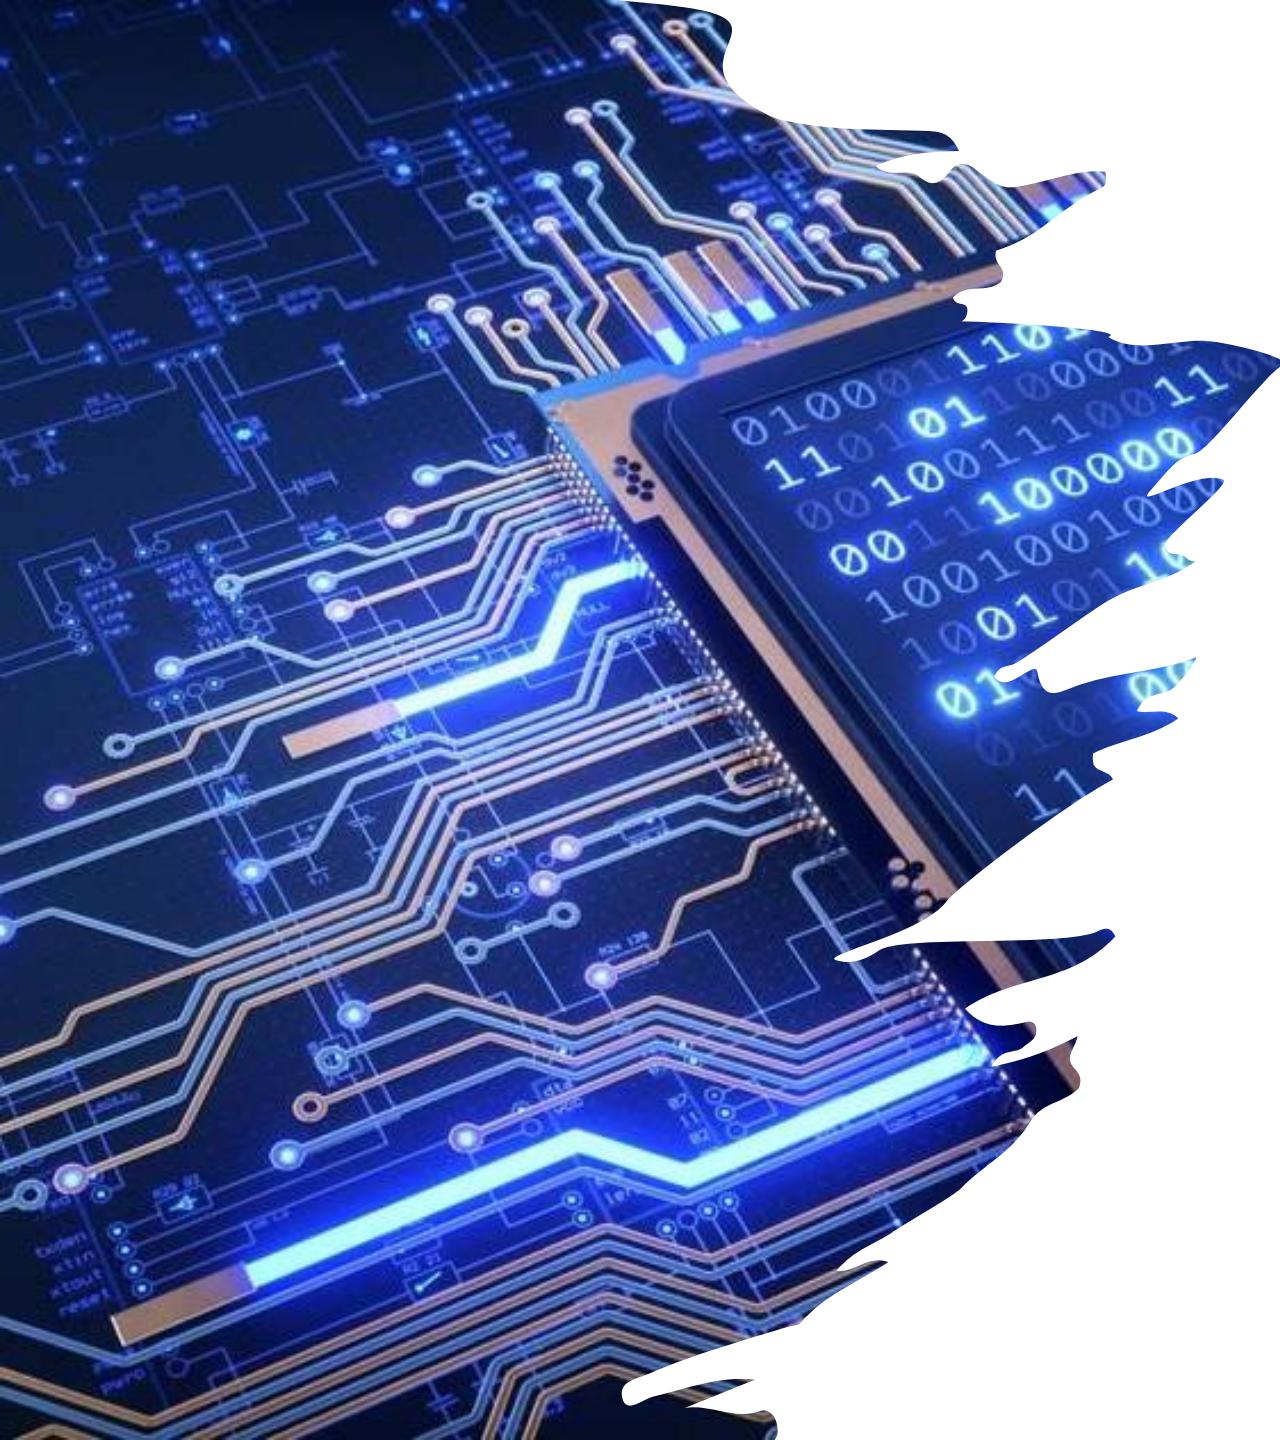
*10-100 times faster
than RAM!*



How does it work?

Accessing a chunk of data in RAM:

- OS will find entry point in RAM
This will require a lookup in virtual memory.
- OS will copy data to a memory cache
If entire chunk of data is too large, only parts will be copied.
This is a slow and daunting task.
- Our program can now read and write to copied data.
- OS will update RAM when it is ready

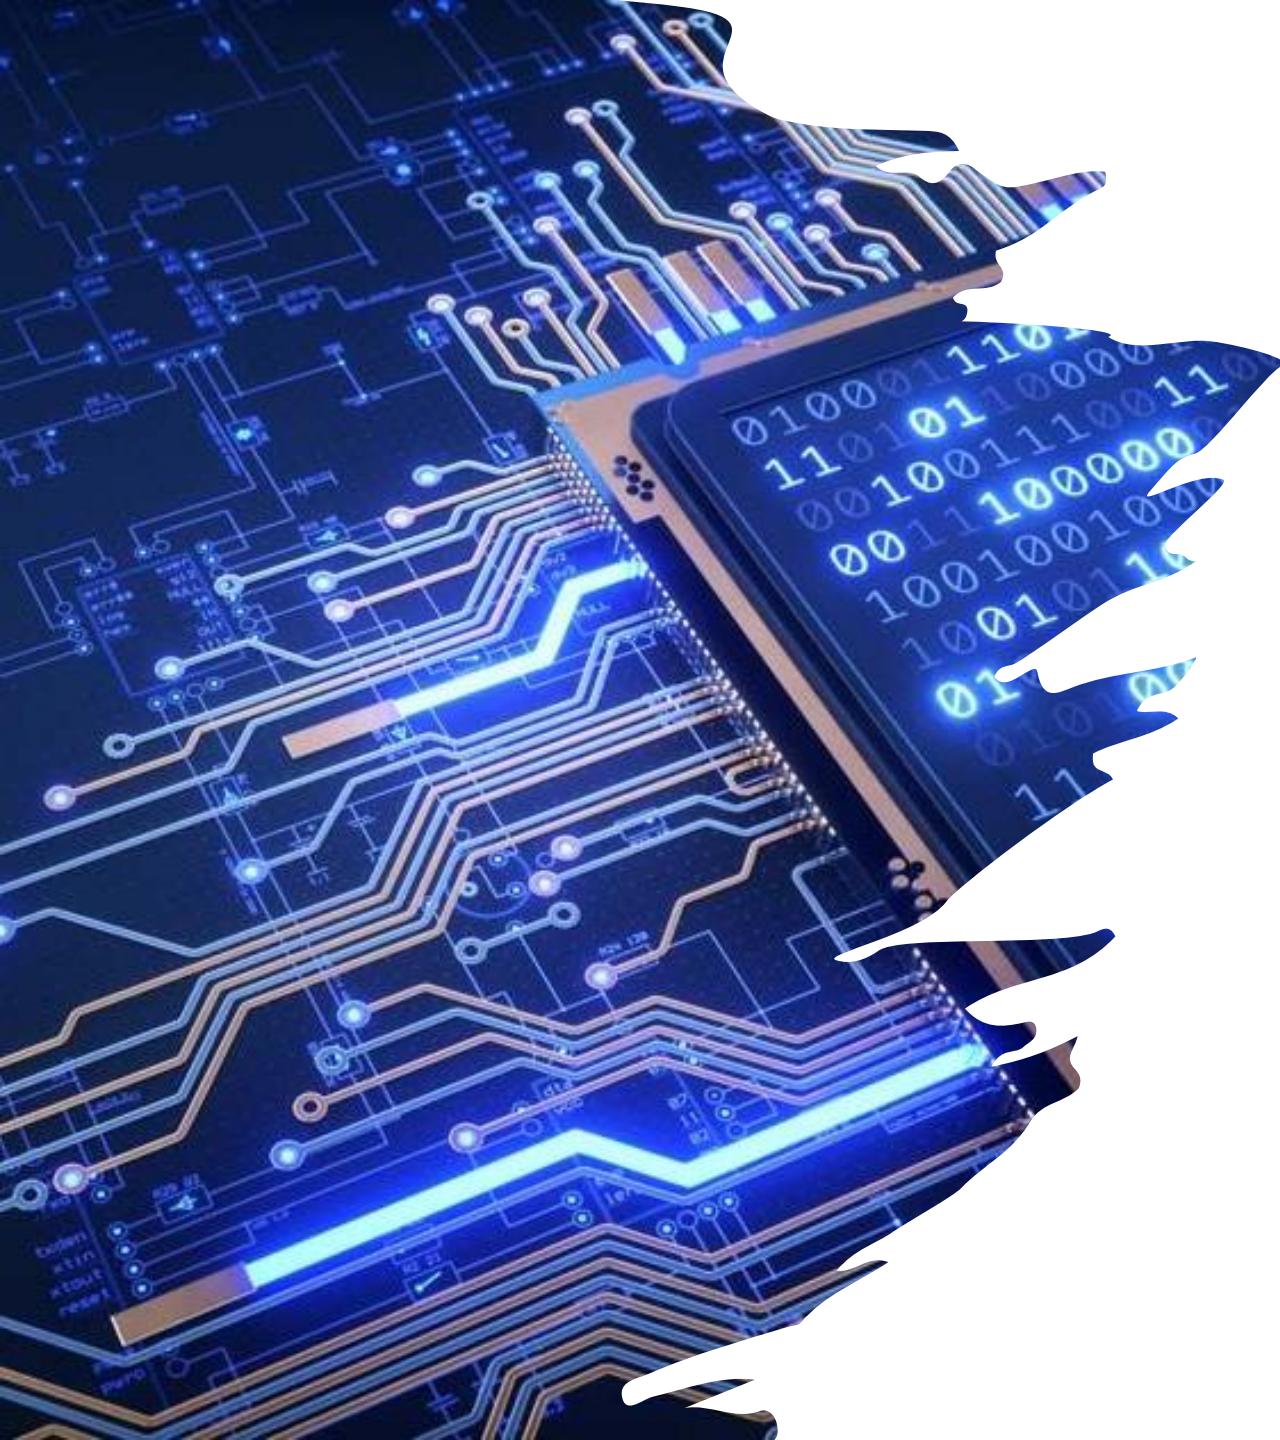


Fast code but how?

```
int [,] numbers = { { 1, 4, 2 }, { 3, 6, 8 } };

for (int i = 0; i < numbers.GetLength(0); i++)
{
    for (int j = 0; j < numbers.GetLength(1); j++)
    {
        Console.WriteLine(numbers[i, j]);
    }
}
```

How many arrays are there in this example?



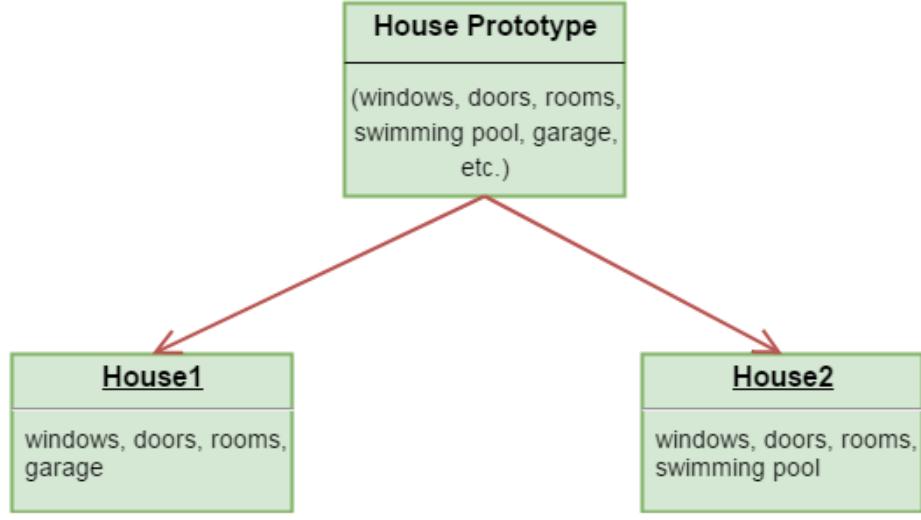
Fast code but how?

```
int [] numbers = { 1, 4, 2, 3, 6, 8 };

for (int i = 0; i < numbers.GetLength(0); i++)
{
    Console.WriteLine(numbers[i]);
}
```

Less code, easier to read, better performance!

OOP, is it any good?



- Pros:
 - Easy to work with
 - Easy to visualize
 - Easy to architect
 - Easy to maintain
 - ...

OOP, is it any good?

Example:

- We have created a city building game
- All buildings, including all interior, is defined in a complex hierarchy of classes
- Write an algorithm which will visit all buildings to check if they are locked.

Any performance issues?



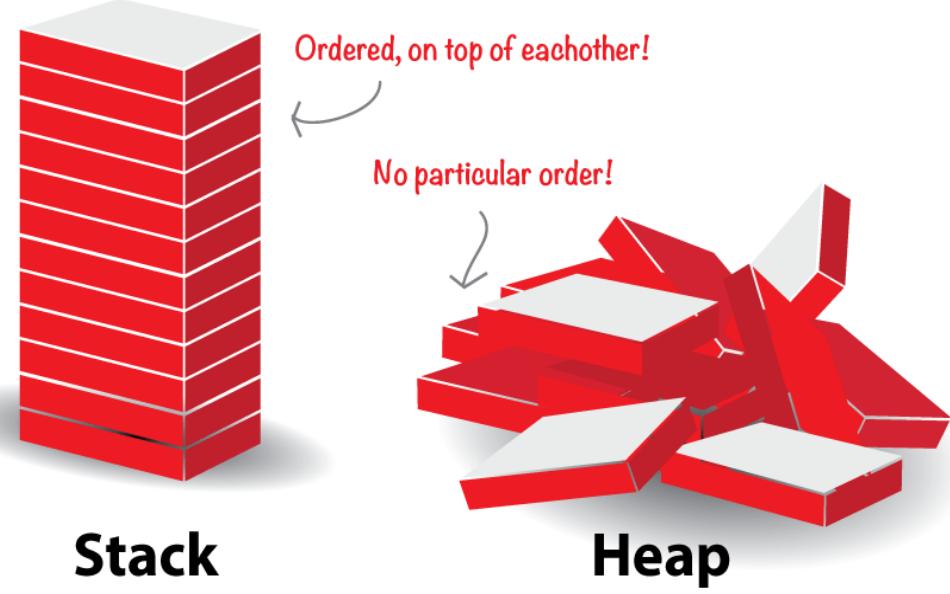
OOP, is it any good?

Answer:

- If a building is built up by multiple classes, we will need to cache a lot of data, slow performance.
- We only need a Boolean to check if a building is accessible or not
- If only we had one array of Booleans for all buildings to check if they were locked or not

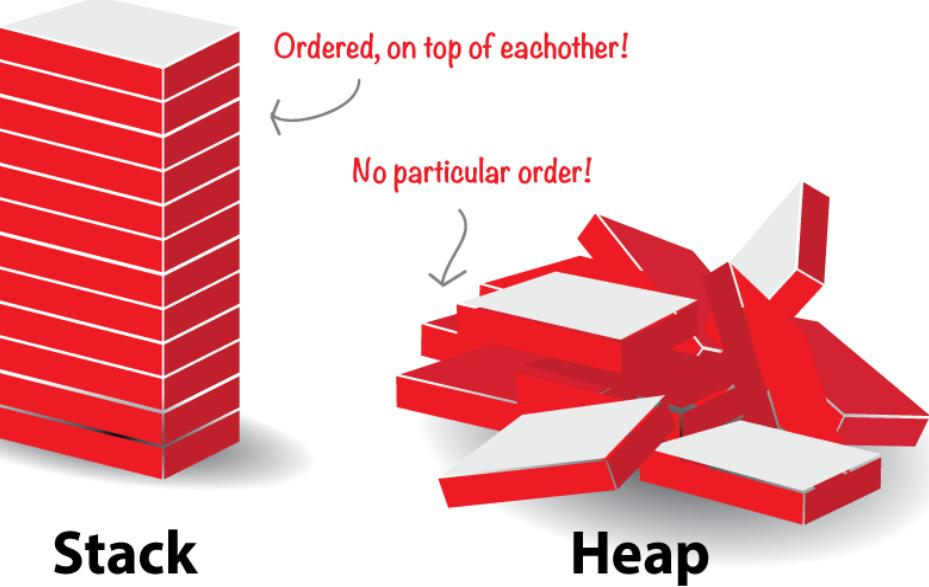


Enable Data Driven Design for Unity!



Stack vs Heap

WHY DO WE NEED TO KNOW
THE DIFFERENCE?



Stack vs Heap

WHY DO WE NEED TO KNOW
THE DIFFERENCE?

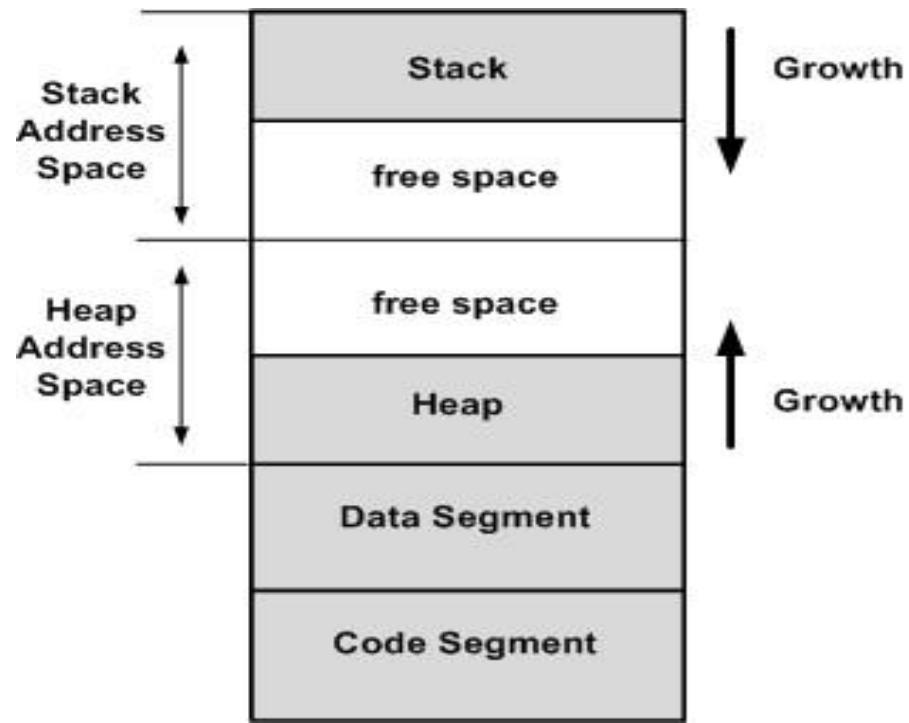
Because we use memory all the time!

Stack

- Local variables
- Memory is stacked (ordered)
- Fast memory allocation
- Automatically removed
- No risk of memory leak
- Limited memory size

Heap

- Global variables
- Memory stored randomly?
- Slow memory allocation
- Not removed automatically
- Risk of memory leak
- Unlimited memory size



Stack vs Heap

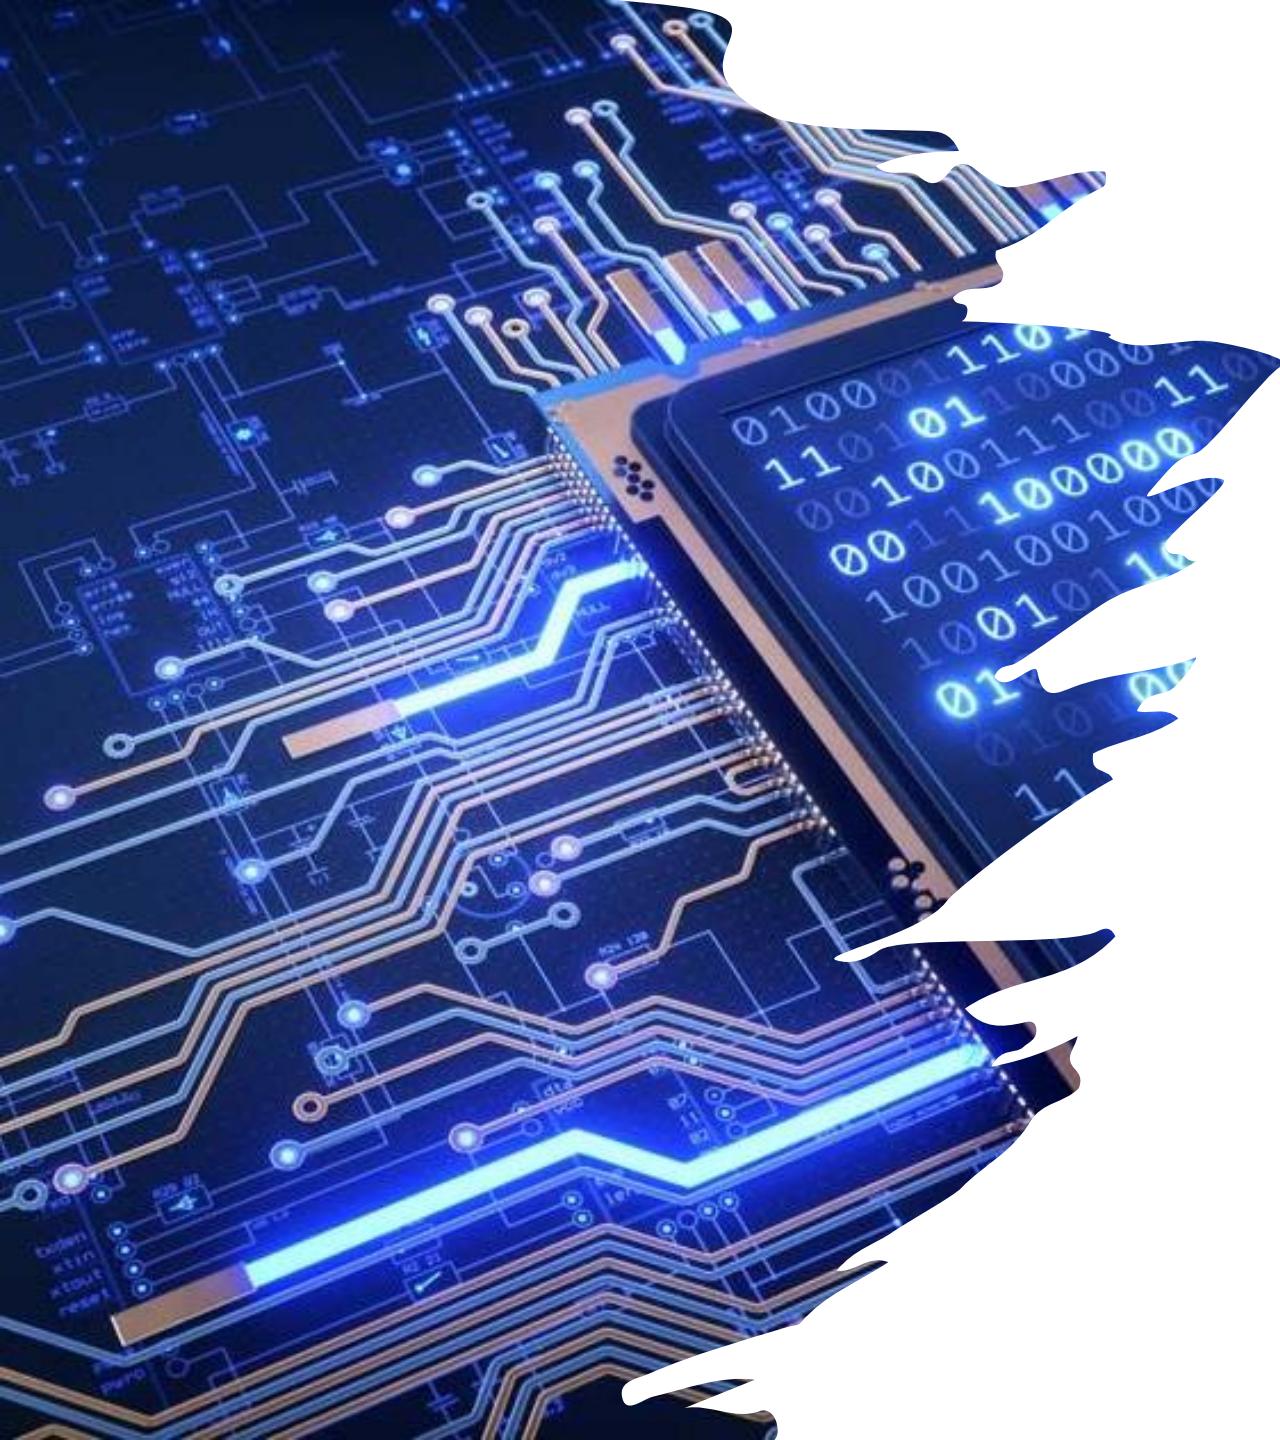
ALWAYS USE **STACK** IF YOU CAN
GET AWAY WITH IT... IT'S FAST!

**Knowledge of
how computer
memory works is
key for fast code!**





**A fast CPU will not help
you, if you don't have
an efficient memory
management!**

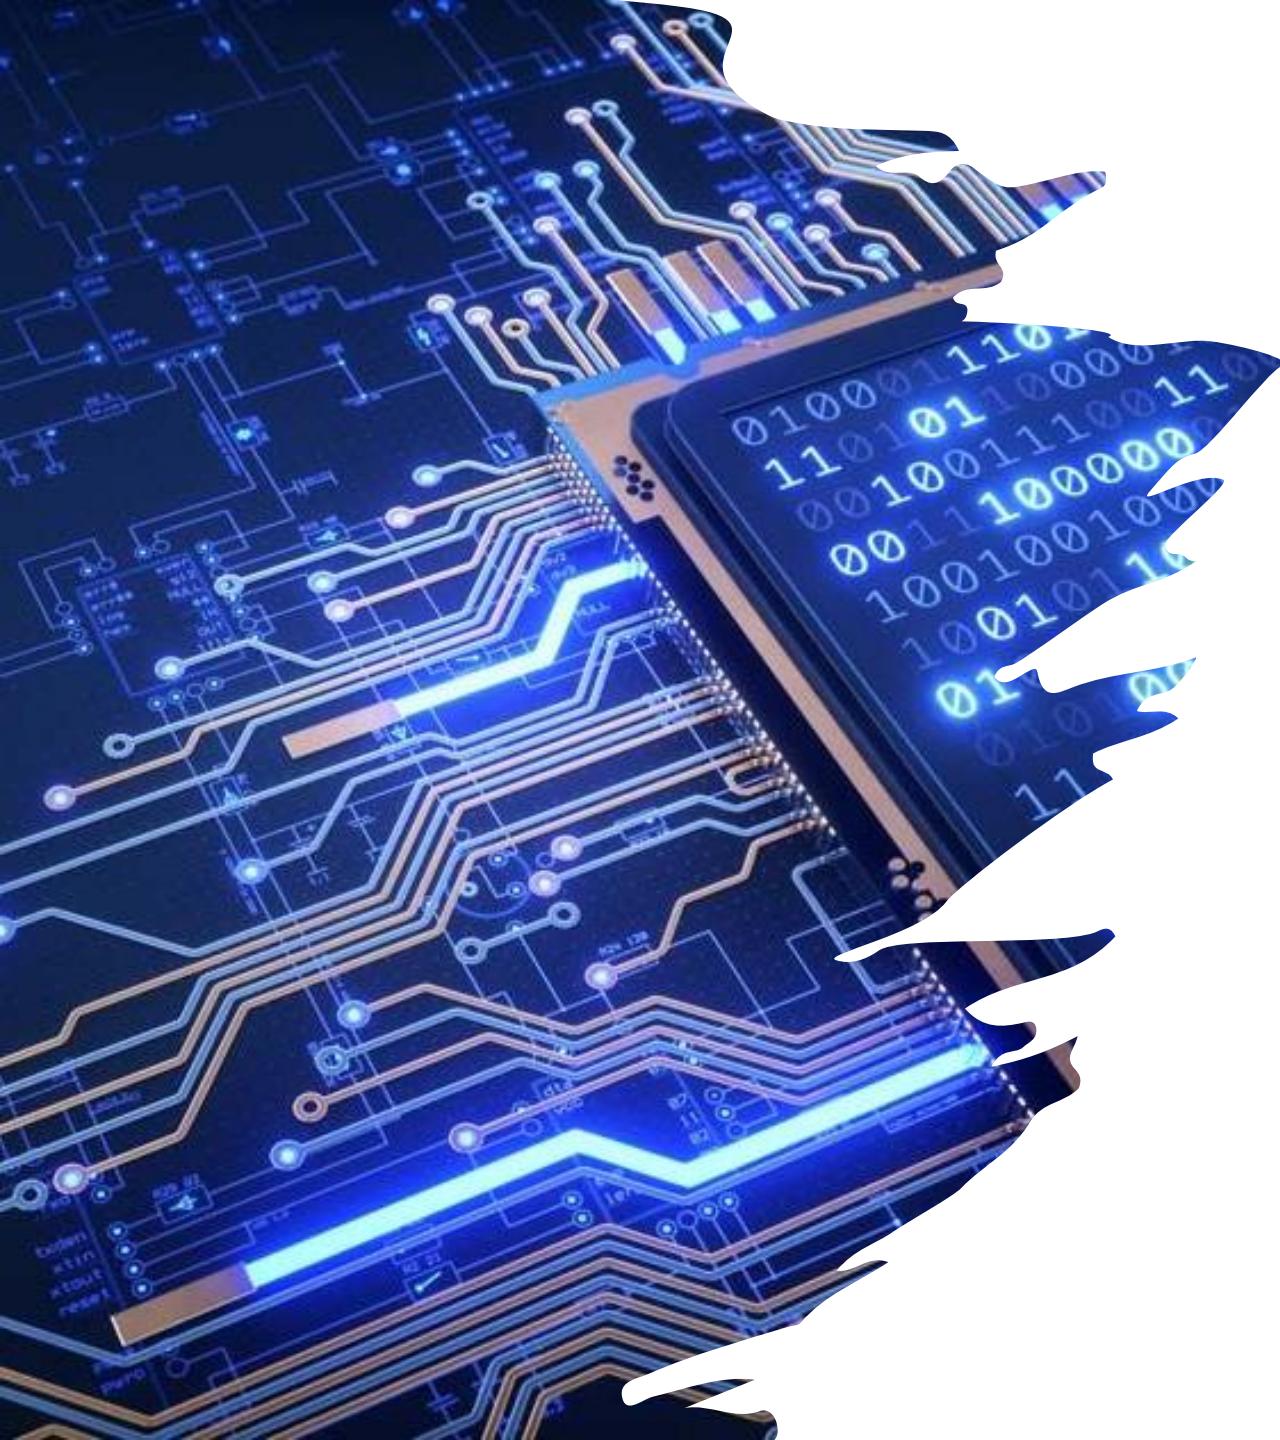


What memory do we use?

```
int [] numbers = { 1, 4, 2, 3, 6, 8 };

for (int i = 0; i < numbers.GetLength(0); i++)
{
    Console.WriteLine(numbers[i]);
}
```

Stack or Heap?

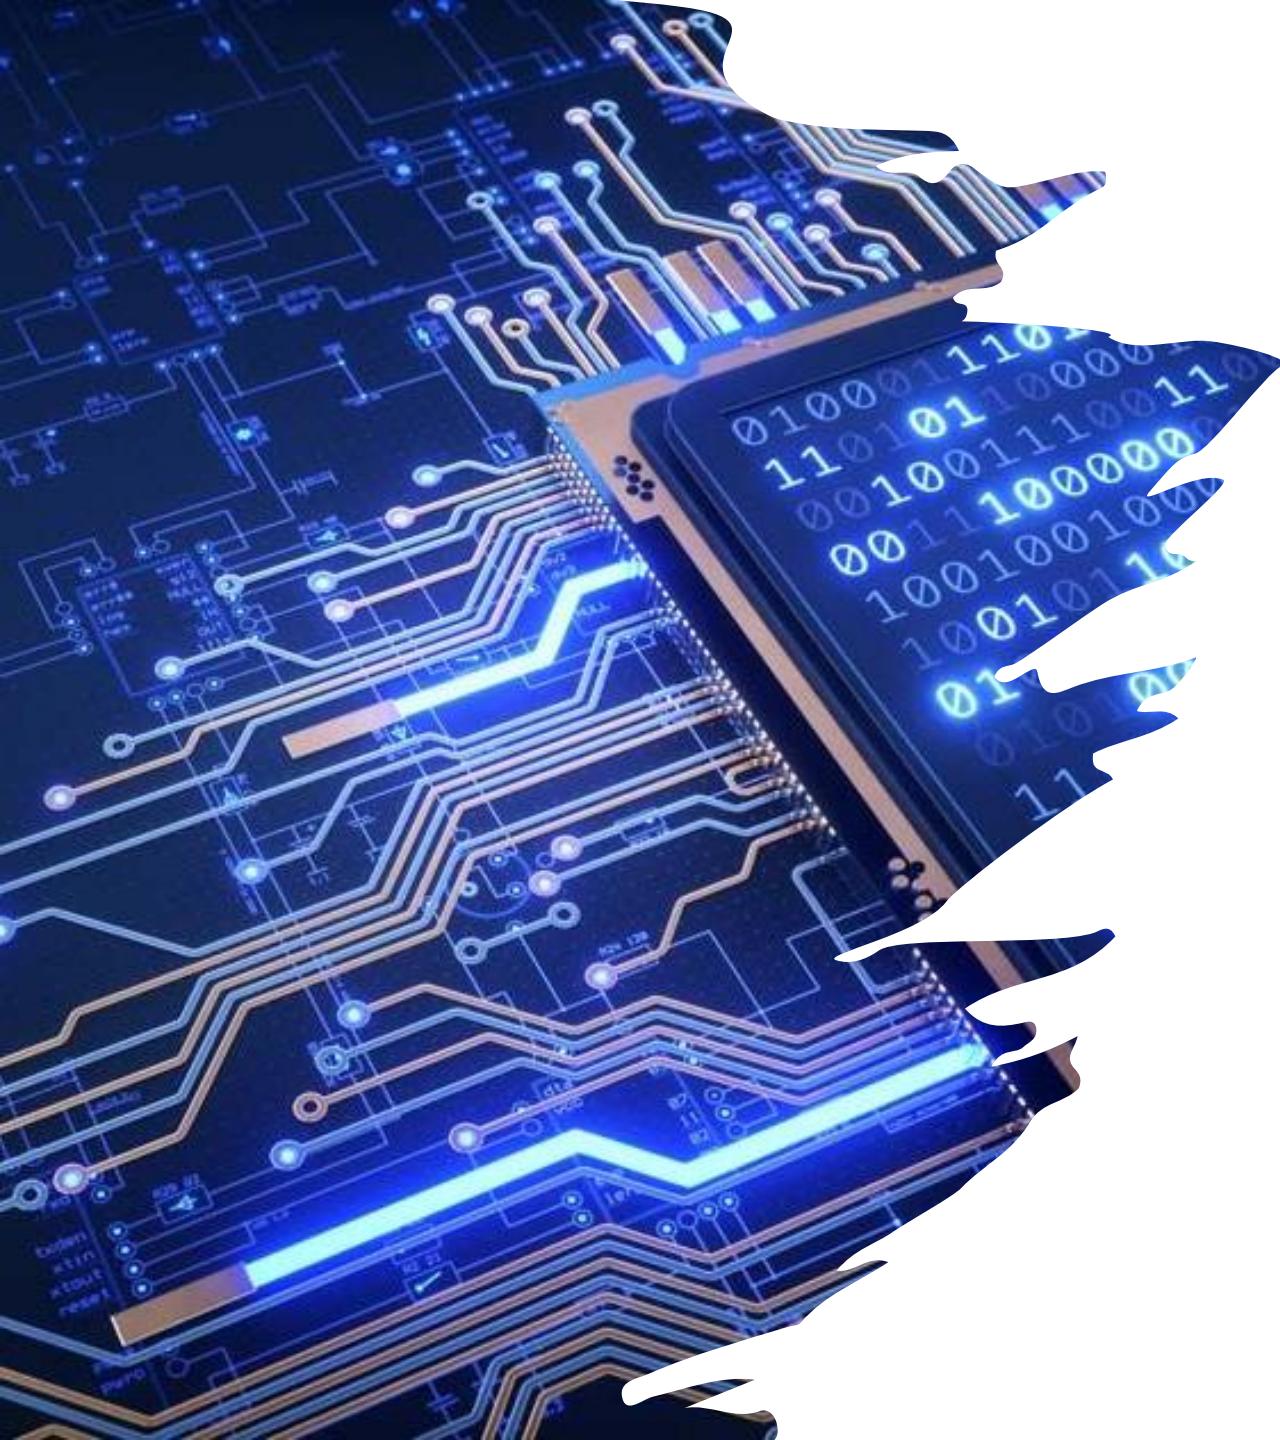


How do we use heap memory?

```
int [] numbers = { 1, 4, 2, 3, 6, 8 };

for (int i = 0; i < numbers.GetLength(0); i++)
{
    Console.WriteLine(numbers[i]);
}
```

Why would we want to use heap memory?

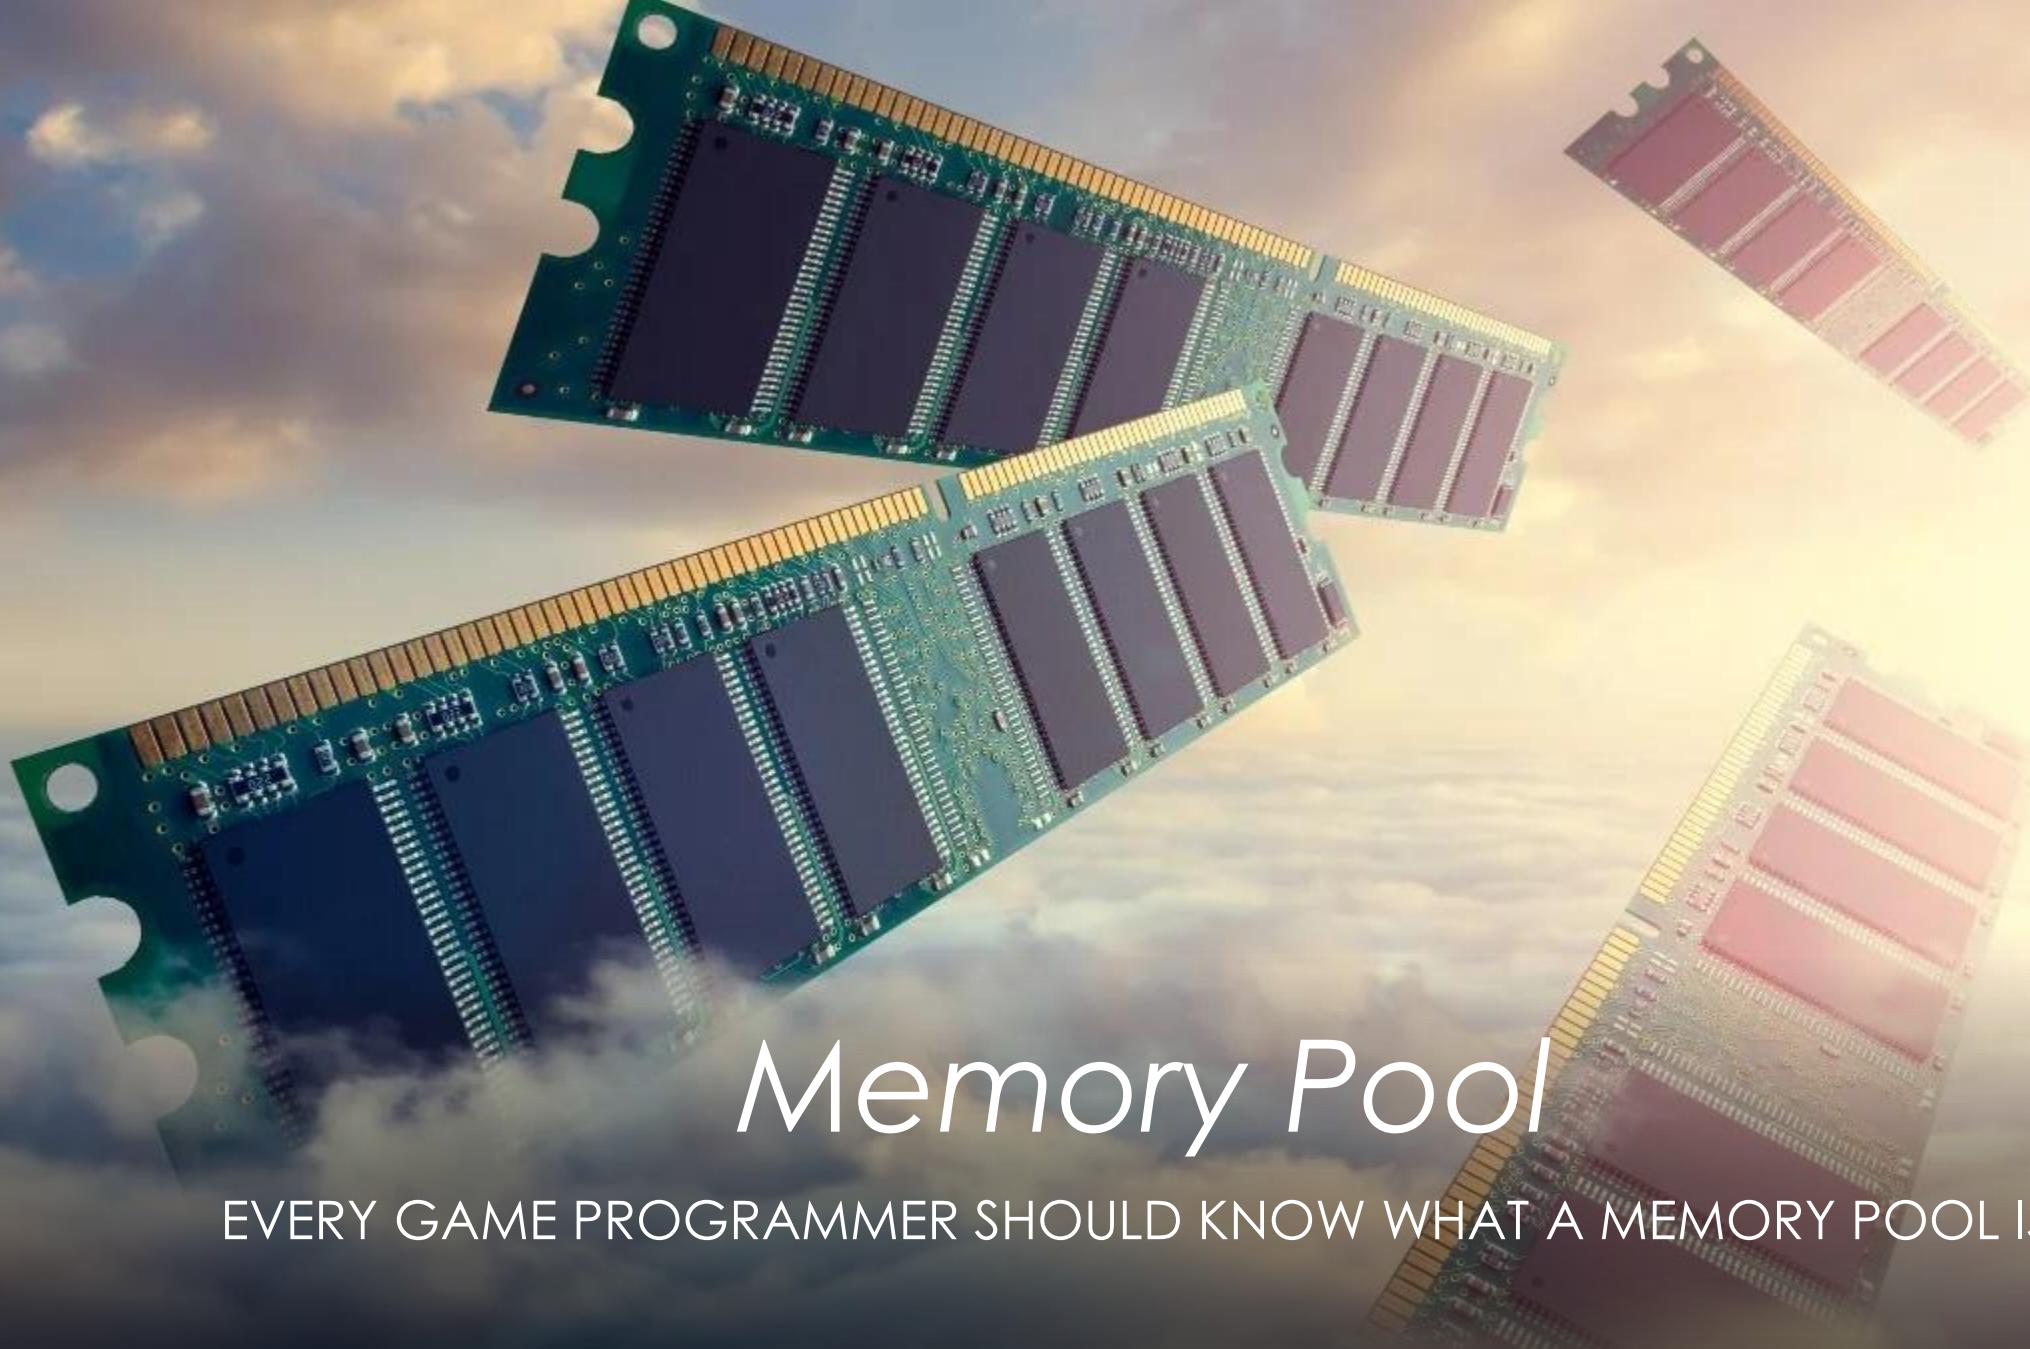


How do we use heap memory?

```
int [] numbers = new int [6];
numbers[0] = 1;
numbers[1] = 4;
...
numbers[7] = 8;

for (int i = 0; i < numbers.GetLength(0); i++)
{
    Console.WriteLine(numbers[i]);
}
```

Heap memory can be used as a memory pool!



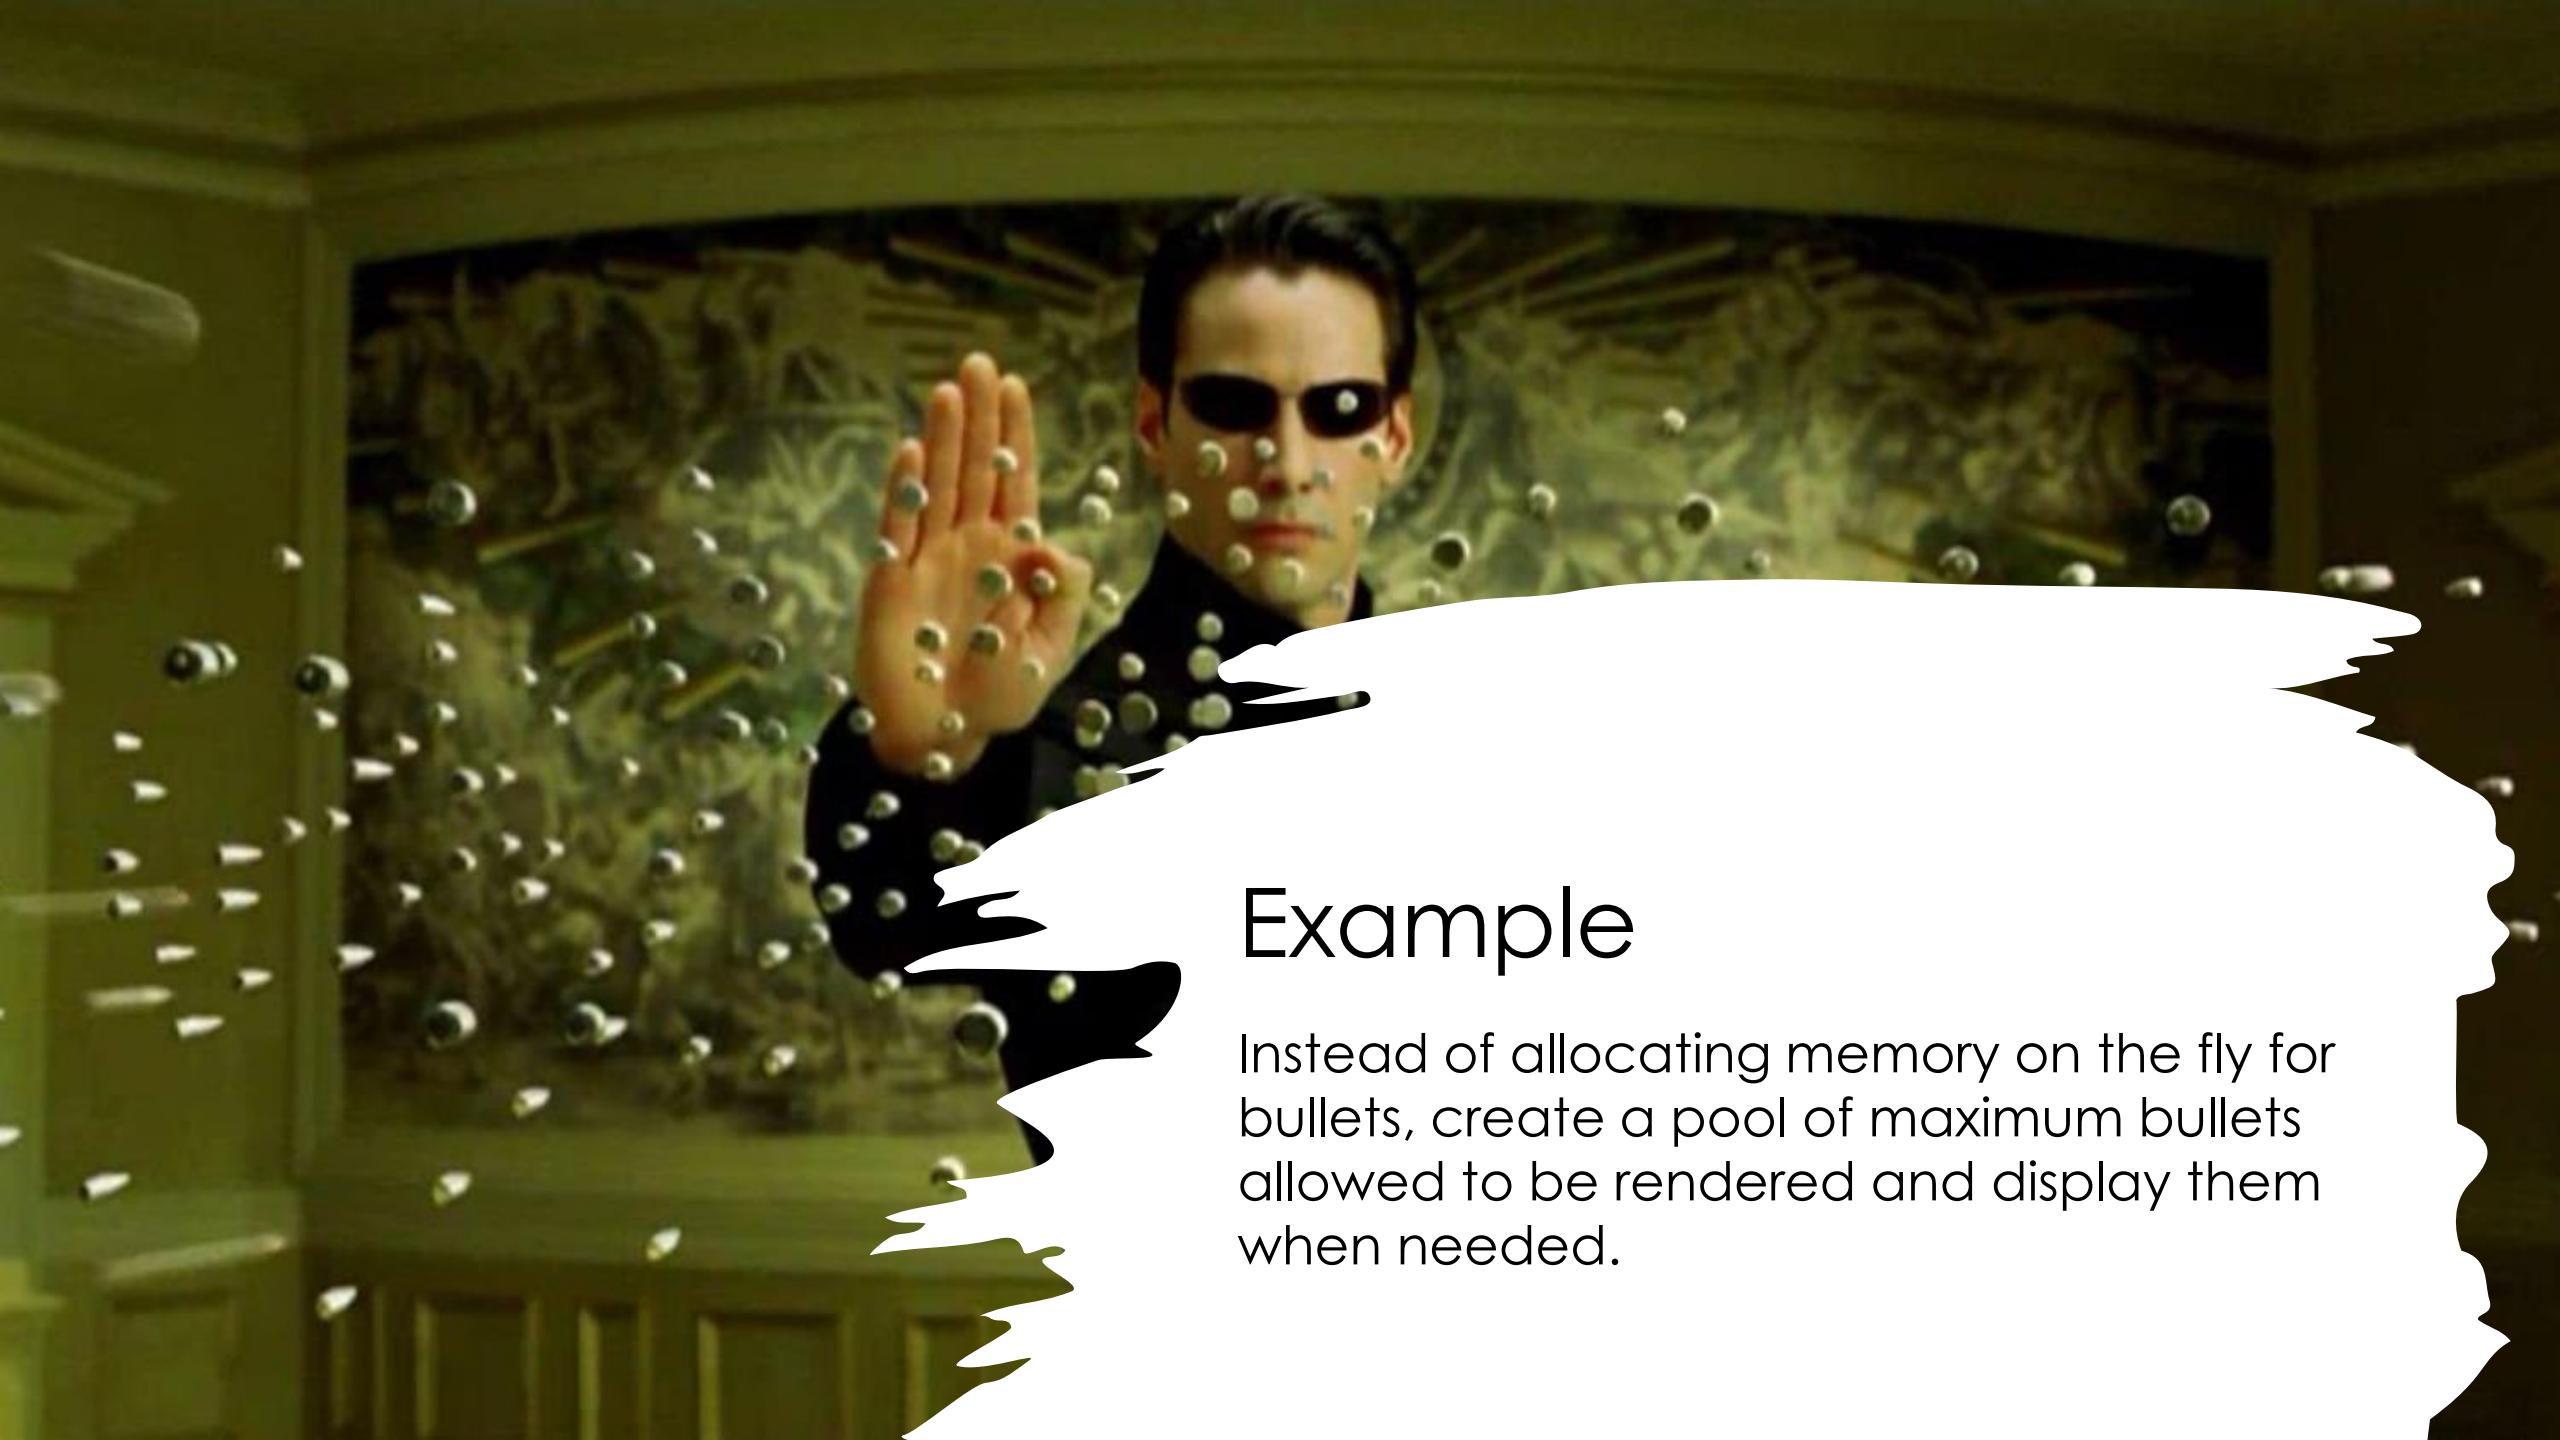
Memory Pool

EVERY GAME PROGRAMMER SHOULD KNOW WHAT A MEMORY POOL IS!

Memory Pool

- It is a common technique used to efficiently allocate and deallocate memory for data structures and objects during program execution.
- Instead of allocating memory separately on the fly, allocate a pool of memory and use it throughout the course of the game.



A scene from the movie The Matrix. Keanu Reeves as Neo is shown from the chest up, wearing his signature black sunglasses and black leather jacket. He is positioned in front of a wall that has been completely blown apart by numerous white bullet holes, creating a jagged, torn shape. His right hand is held out towards the viewer, palm facing forward, with several small, glowing green spheres (bullets) visible on his fingers. A white silhouette of a person's head and shoulders is overlaid on the bottom right side of the image.

Example

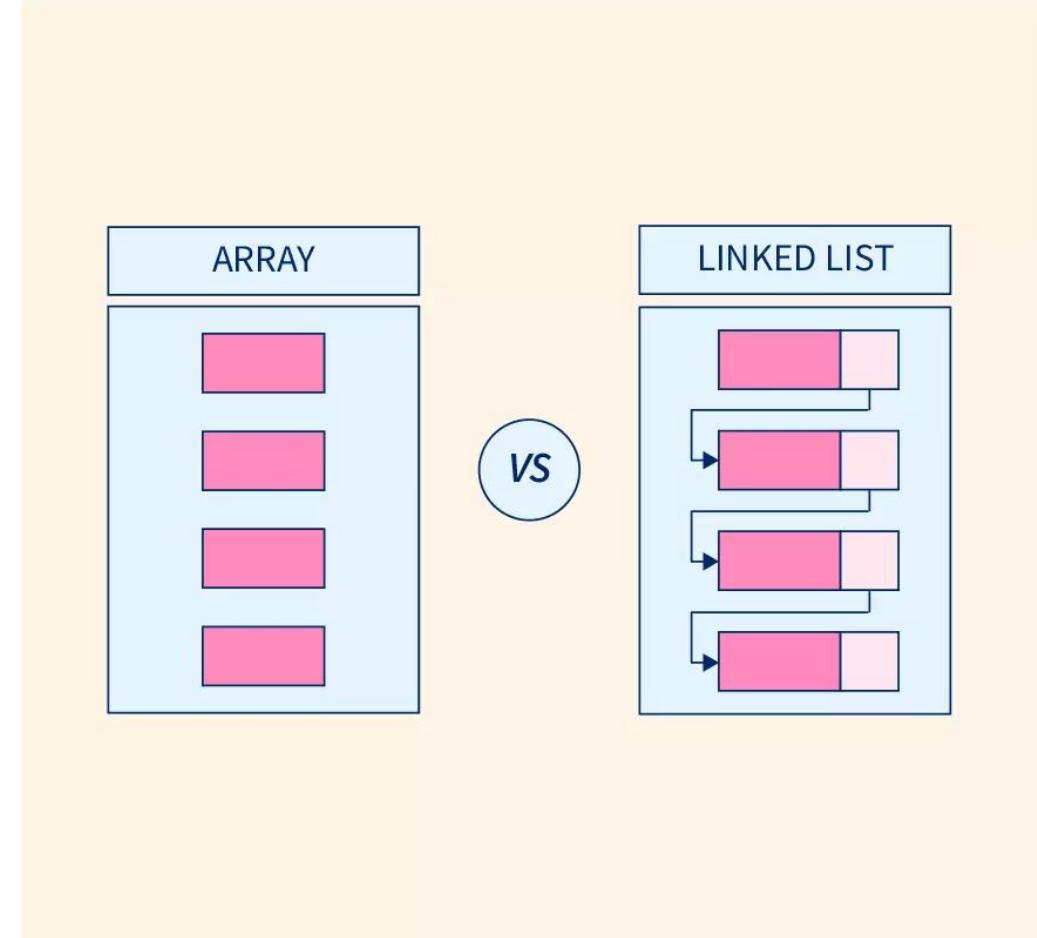
Instead of allocating memory on the fly for bullets, create a pool of maximum bullets allowed to be rendered and display them when needed.

*Most games have
a limited number of
ragdolls available
due to memory
pool*

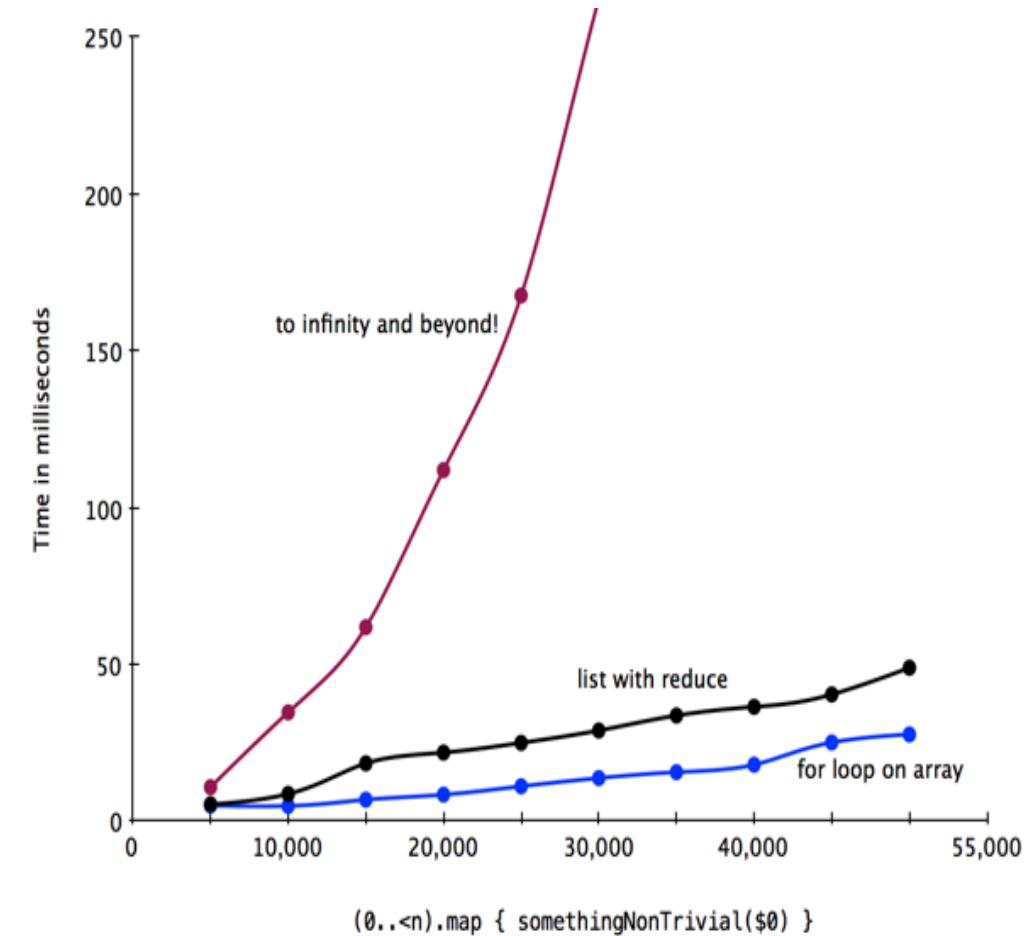


Algorithm & Datastructure

- Key to great performance!



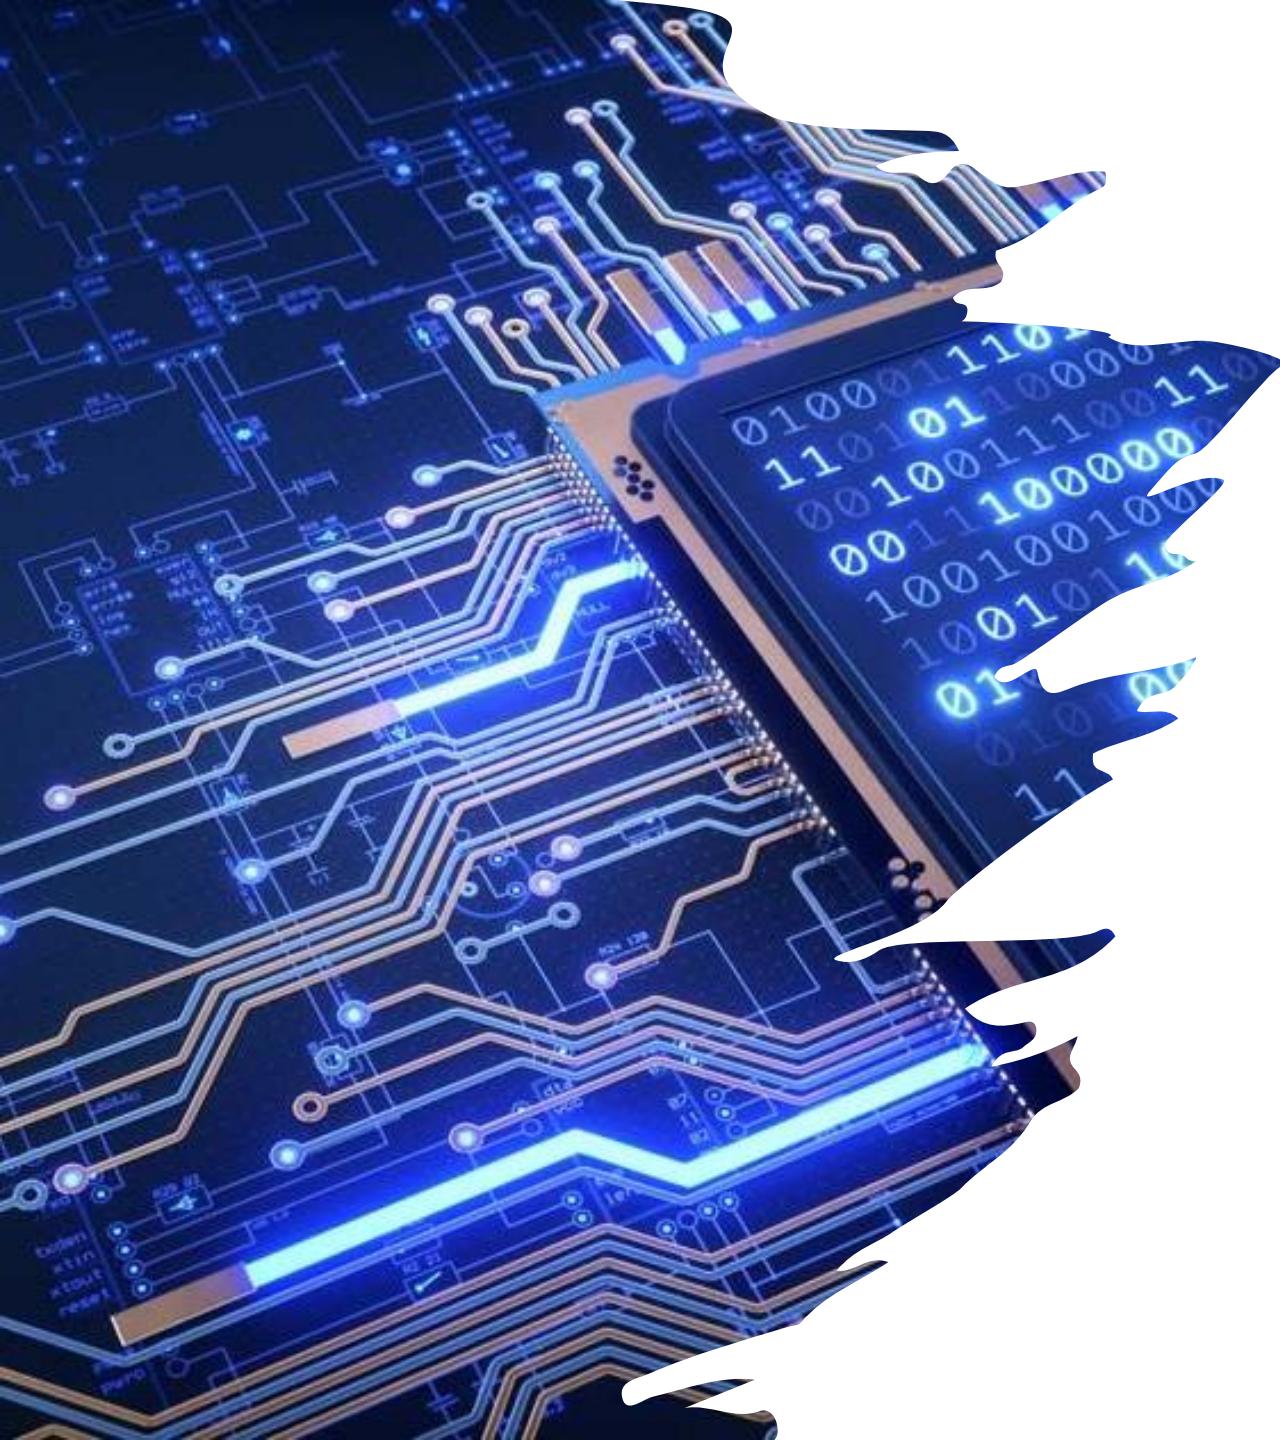
Algorithm & Datastructure



It is important to use the correct tool for the job!



*For fast
performance,
learn to think
like a computer!*



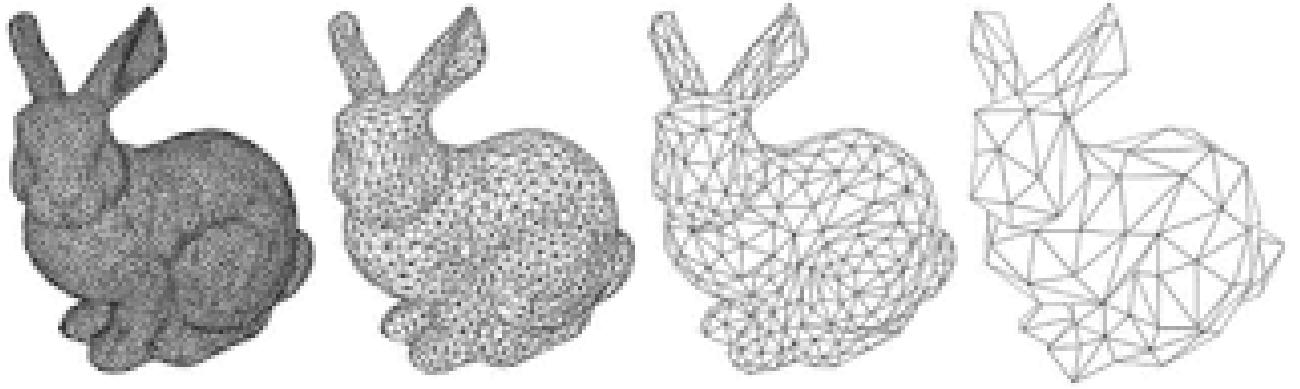
Fast code but how?

```
void foo(bool something)
{
    if (check() && something)
    {
        Console.WriteLine("Something");
    }
}
```

Does the order in an if-statement matter?

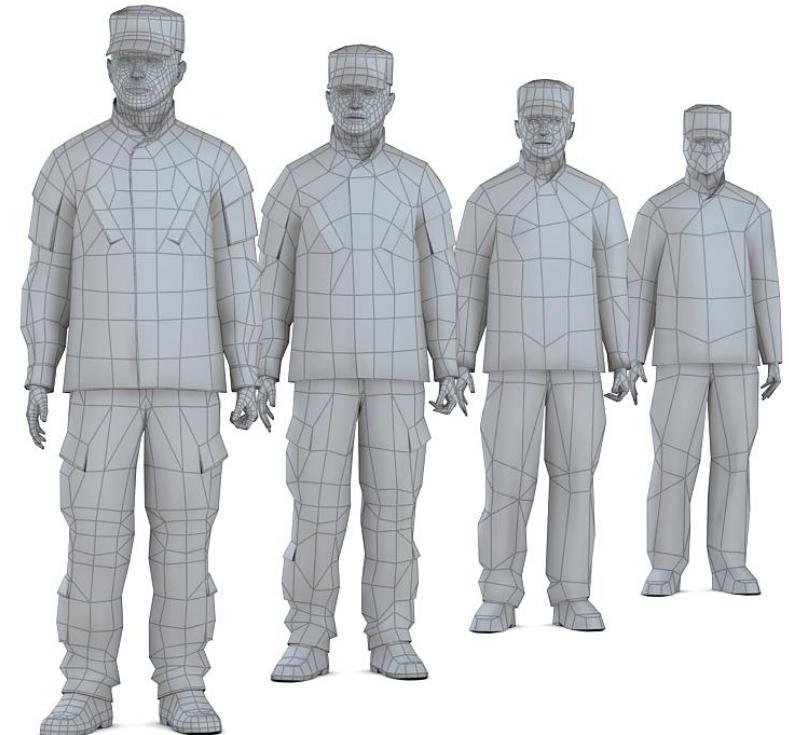
Level of Details (LOD)

WHY SHOULD I USE IT?



*In computer graphics,
LOD refers to the
complexity of 3D
models and textures*

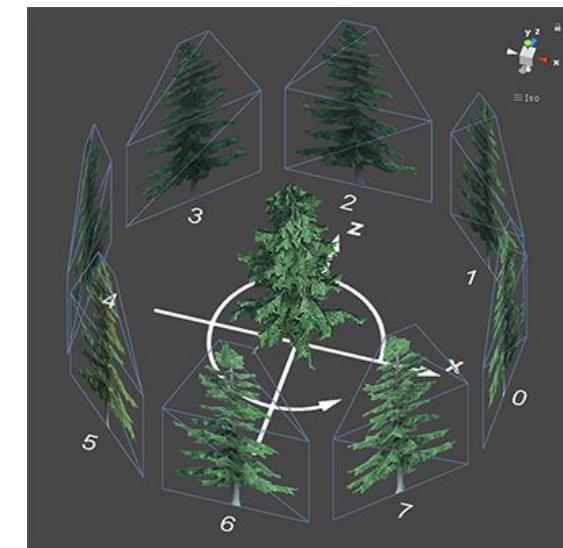
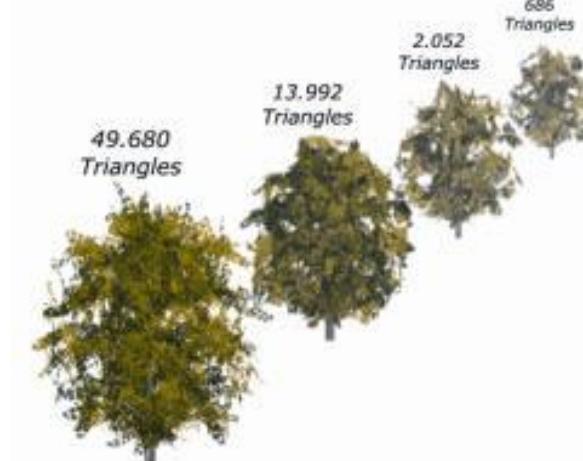
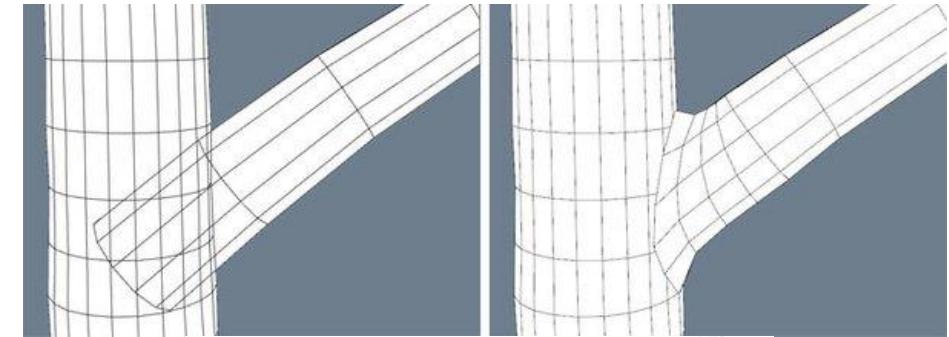
- Details are less important as they get further away from the camera
- LOD can reduce poly-count and texture resolution for objects further away to boost performance
- A perfect system allows a user to customize graphics settings to run on any system



LOD & Billboards

Rendering of trees and vegetation is very performance heavy

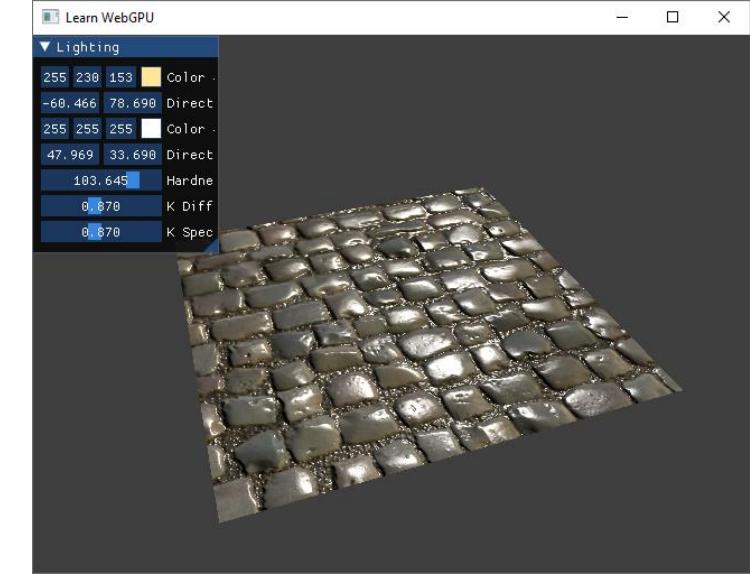
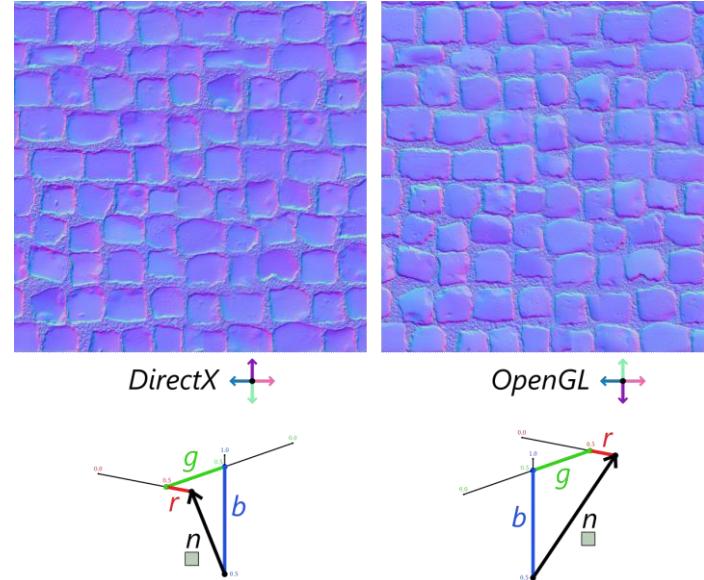
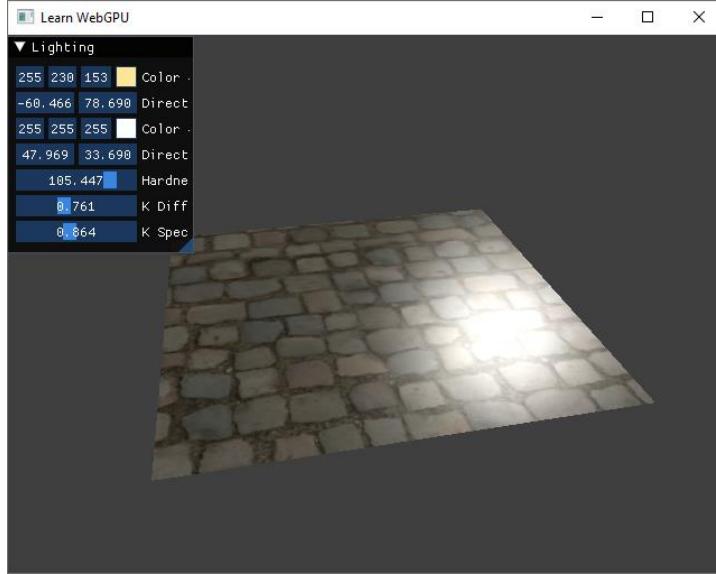
It is a common trick to use LOD, where the lowest level of a vegetation is rendered to a billboard





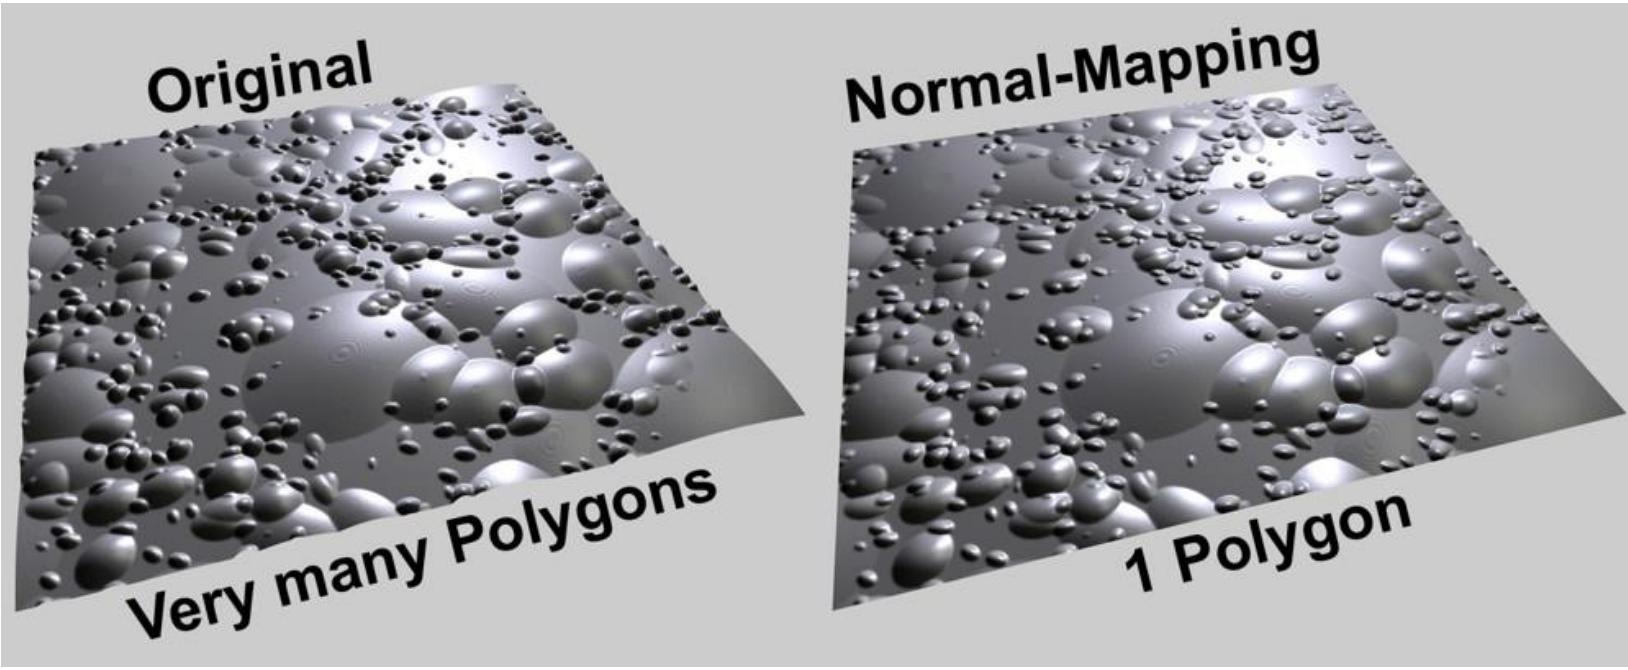
Normal Mapping

WHAT IS IT GOOD FOR?



Normal Mapping

NORMAL MAPPING ADD SURFACE DETAILS
WITHOUT ADDING GEOMETRIC DETAILS



Normal Mapping

NORMAL MAPPING ADD SURFACE DETAILS
WITHOUT ADDING GEOMETRIC DETAILS

Normal Mapping

PERFECT FOR ADDING
DETAILS TO CHARACTERS





Parallax Mapping

- An extension to normal mapping
- Parallax mapping includes height
- Internal shadows can be calculated



Parallax Mapping

GREAT FOR GROUND AND WALL TEXTURES



Parallax Mapping

GREAT FOR ANY BONDED AREAS THAT NATURALLY BORDERING OTHER SURFACES

Render to texture

- Render to texture to imitate reflective surfaces
- **Make sure you do not render more than what is visible in mirror**





Render to texture

- Render to texture to imitate reflective surfaces
- **For unclear reflections like animated water, you can get away with only render very few things**

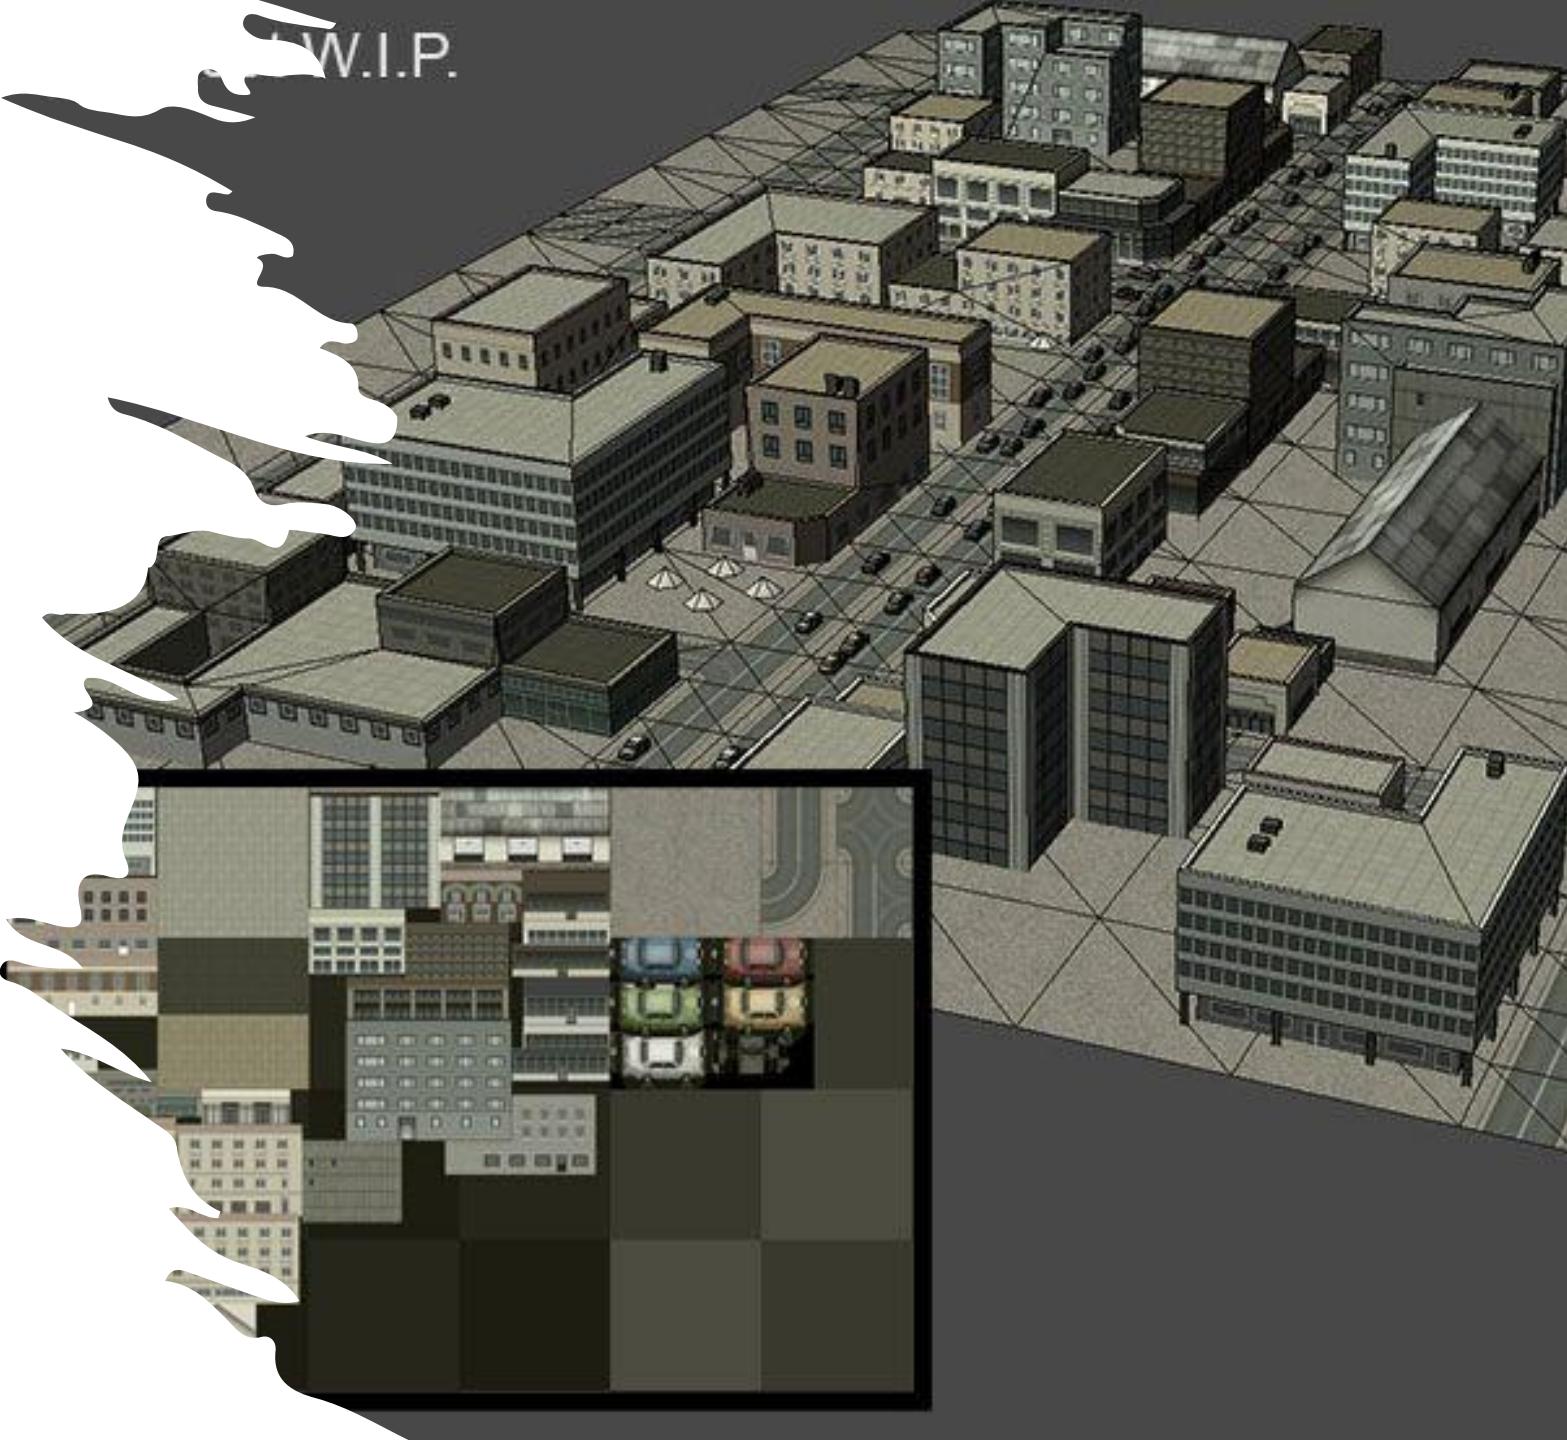


Render to 3D texture

- Rendering buildings in a city is very performance heavy
- Render rooms for buildings, that can be seen from outside, in 3D textures will reduce polygon count significantly

Atlas Textures

- A large texture that has multiple textures combined
- Why is this faster than having many small separate textures?
- What is considered a large texture?
- What does Unity consider to be a large texture?



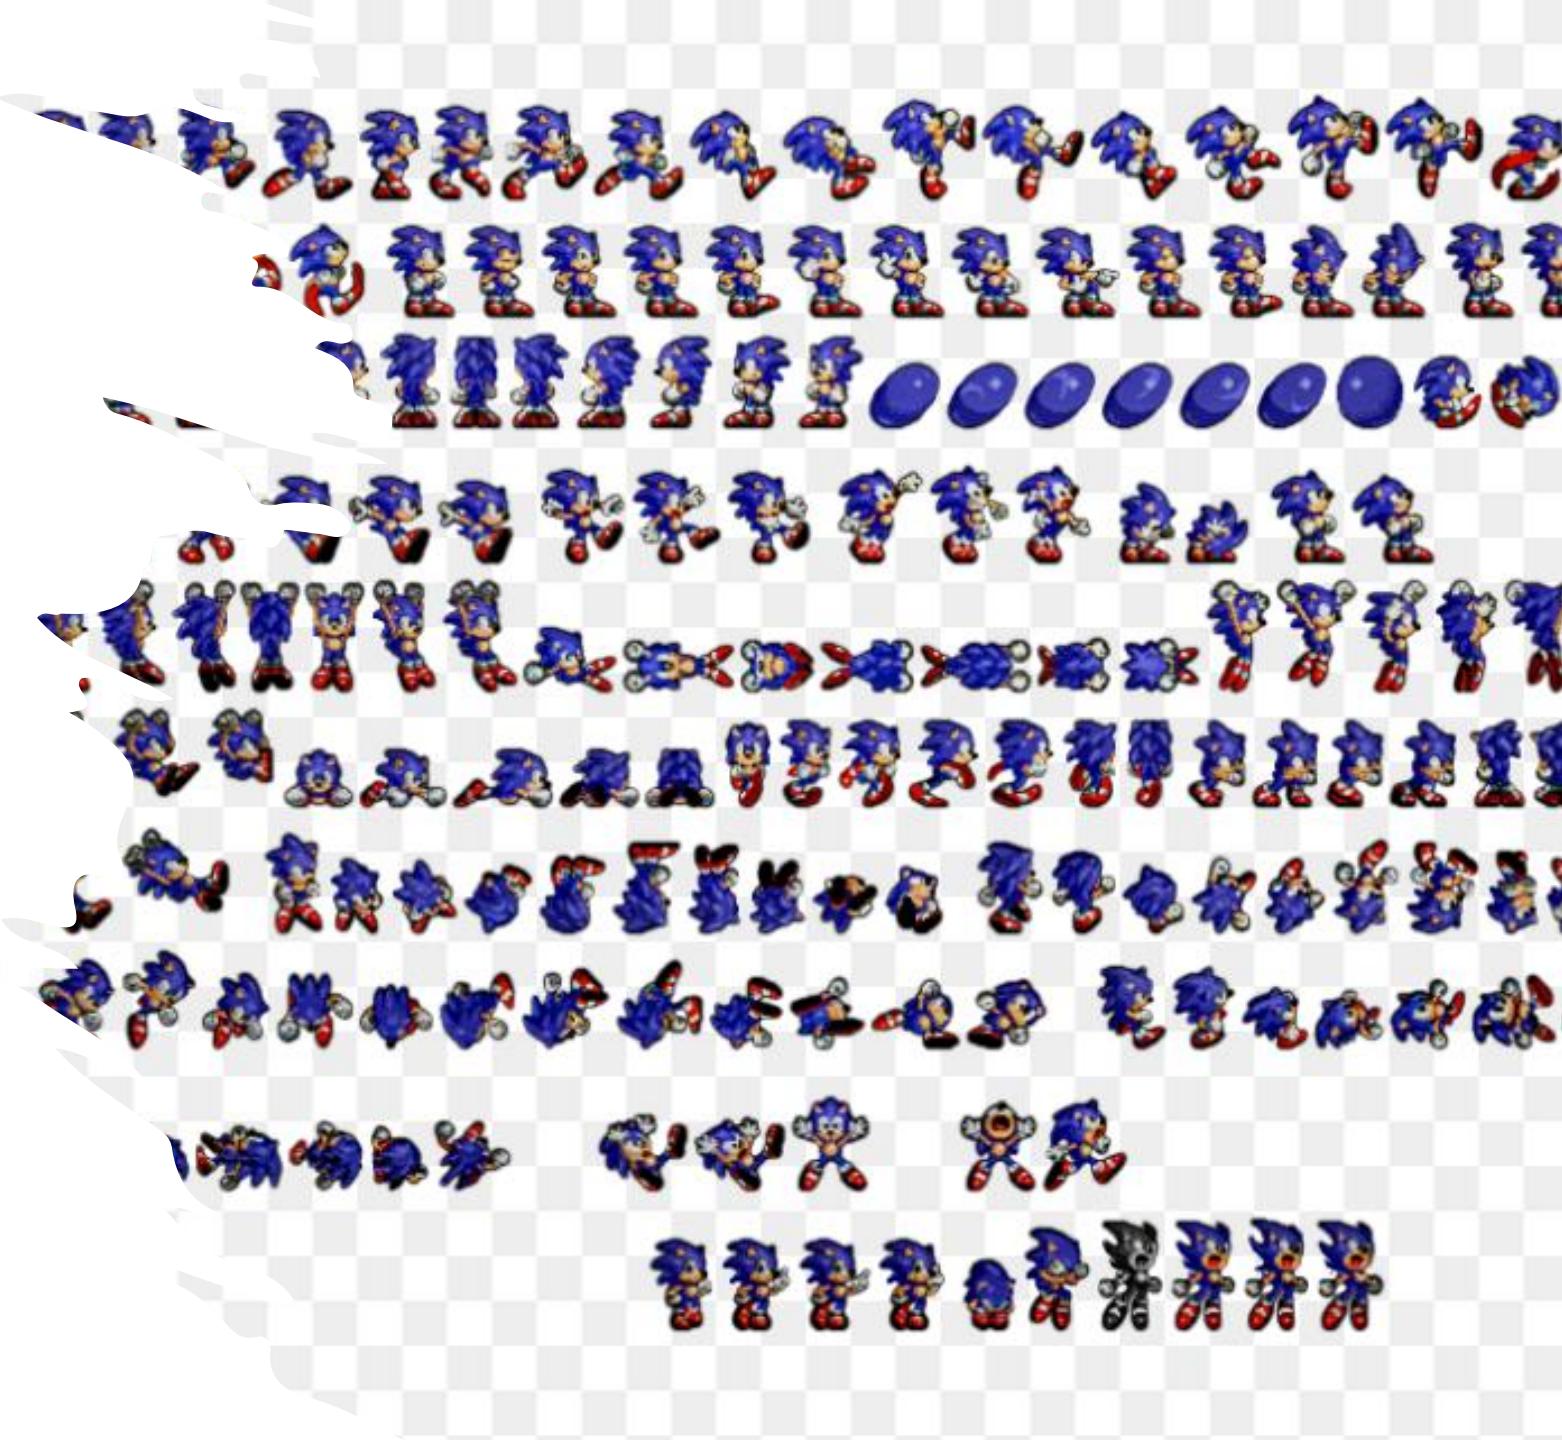
Atlas Texture vs Sprite Sheet

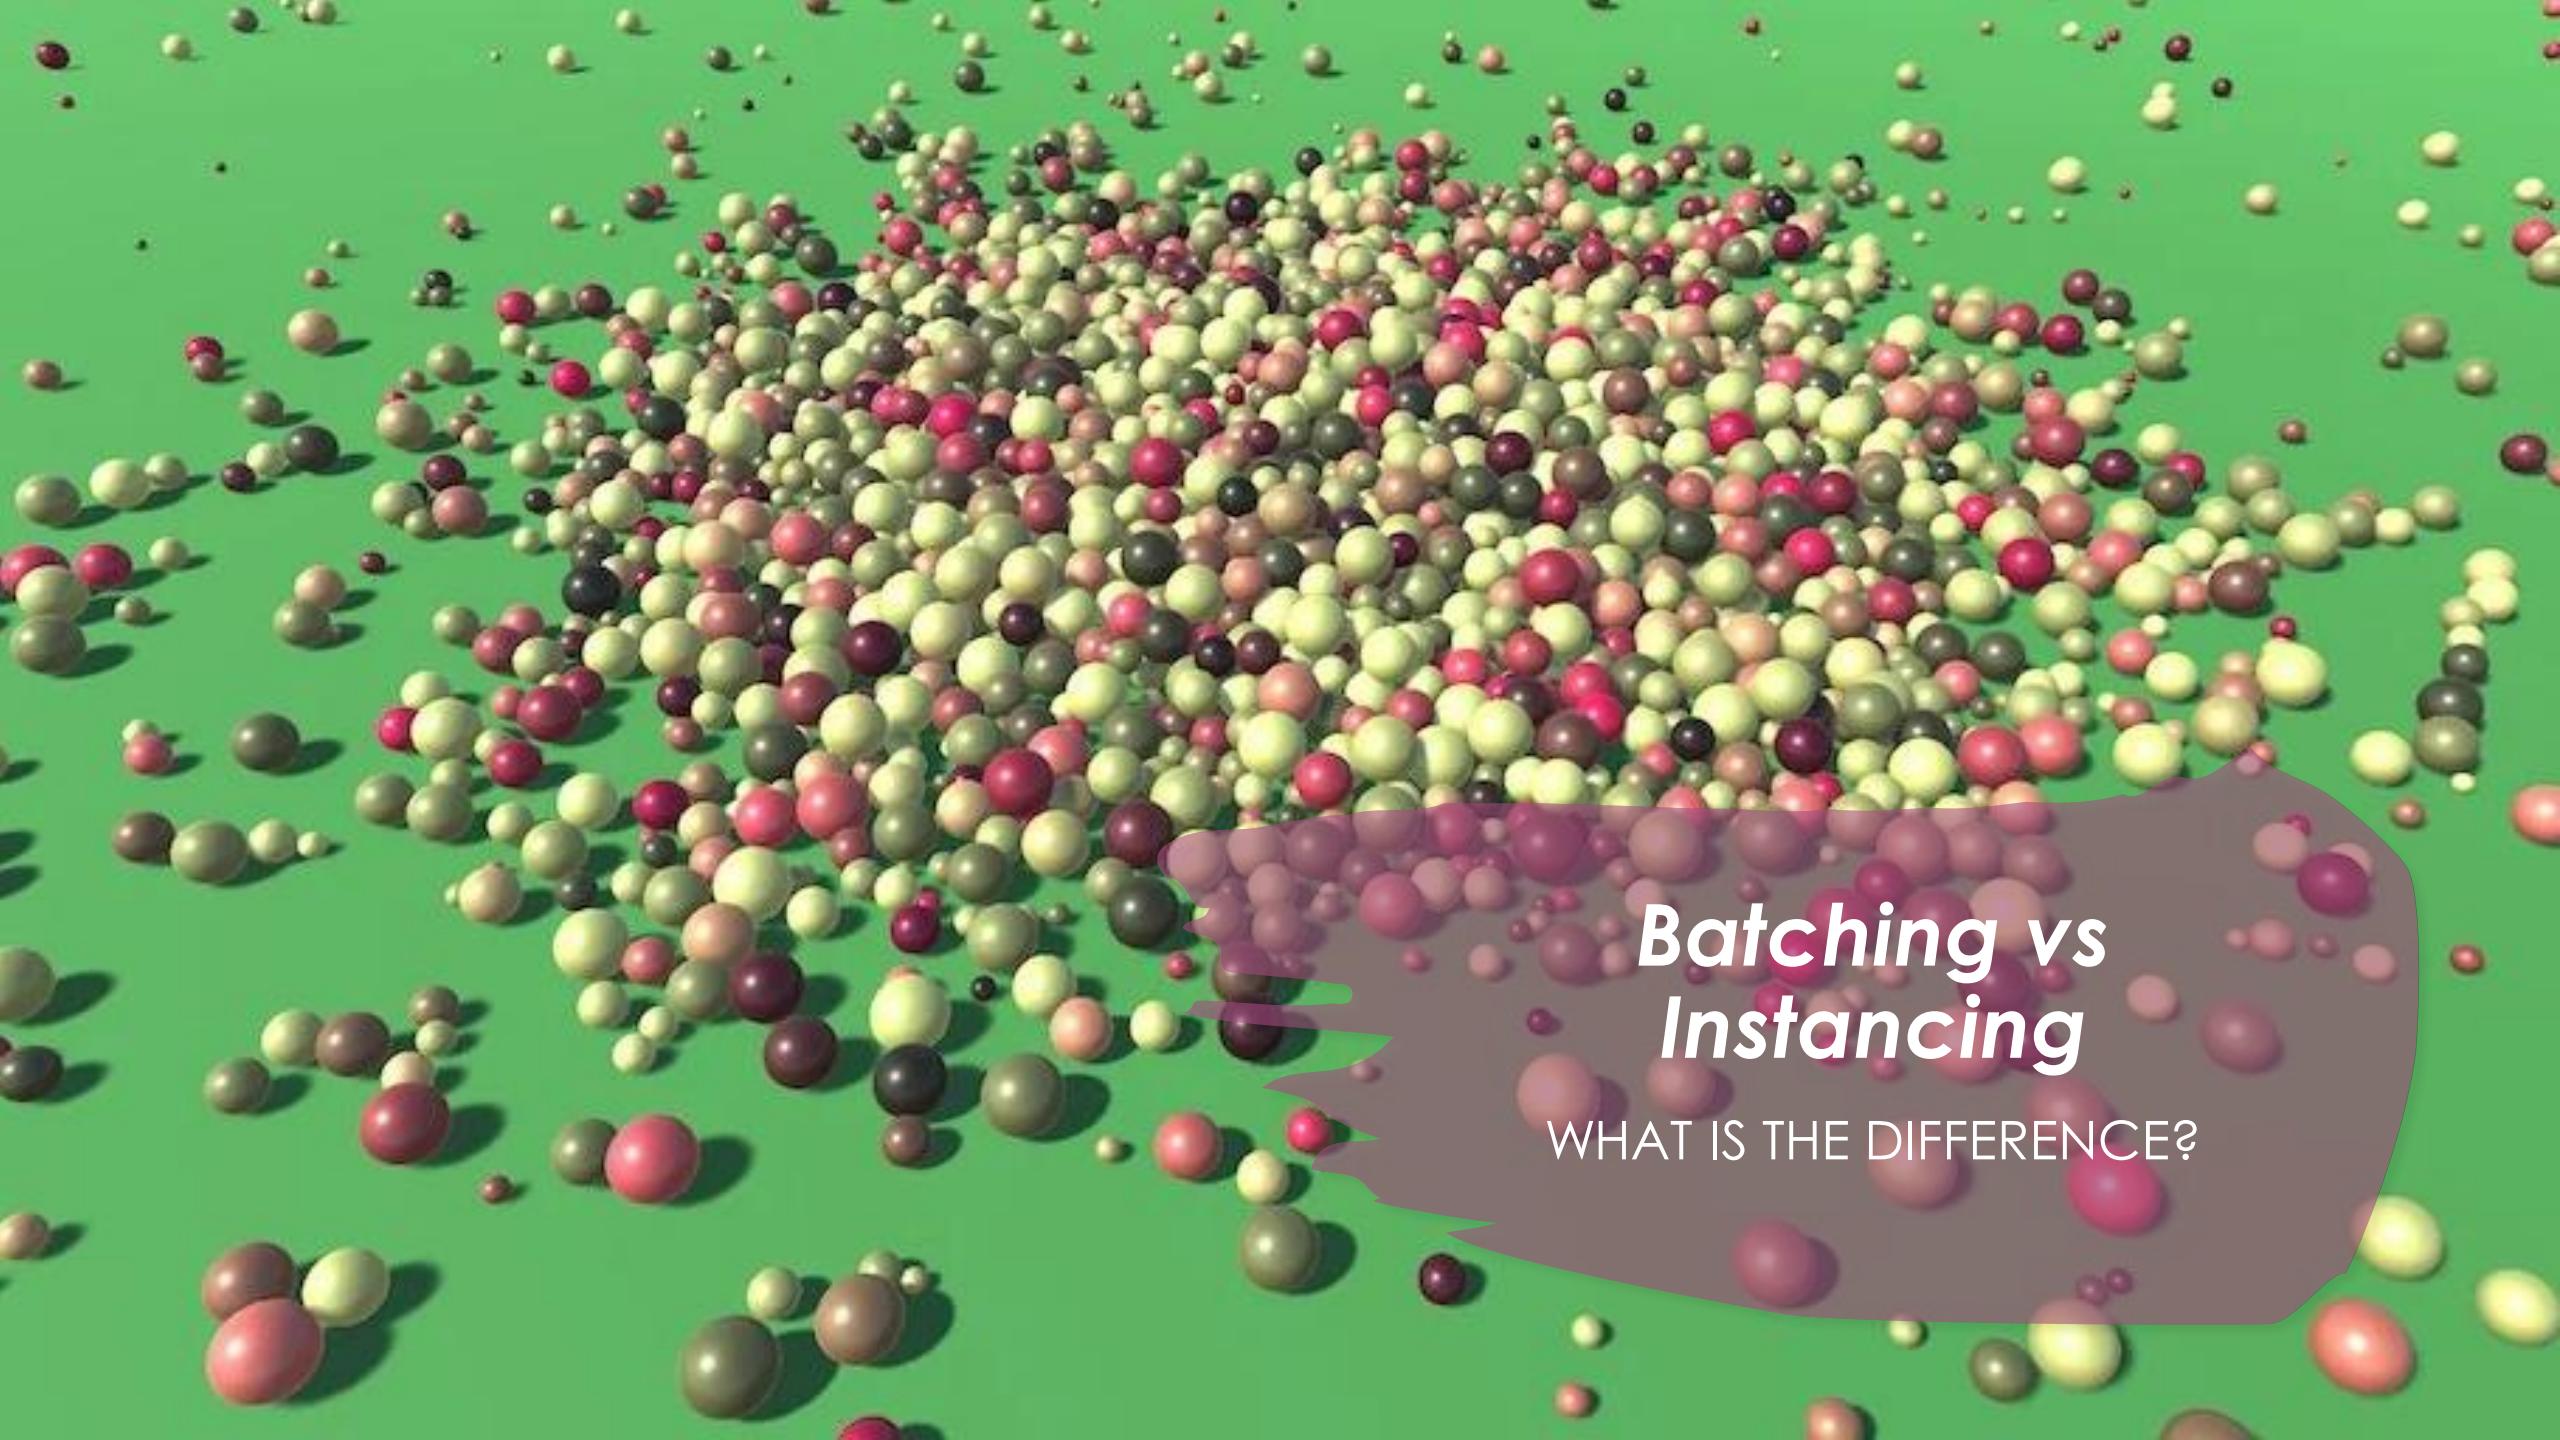
WHAT IS THE DIFFERENCE?



Sprite Atlas

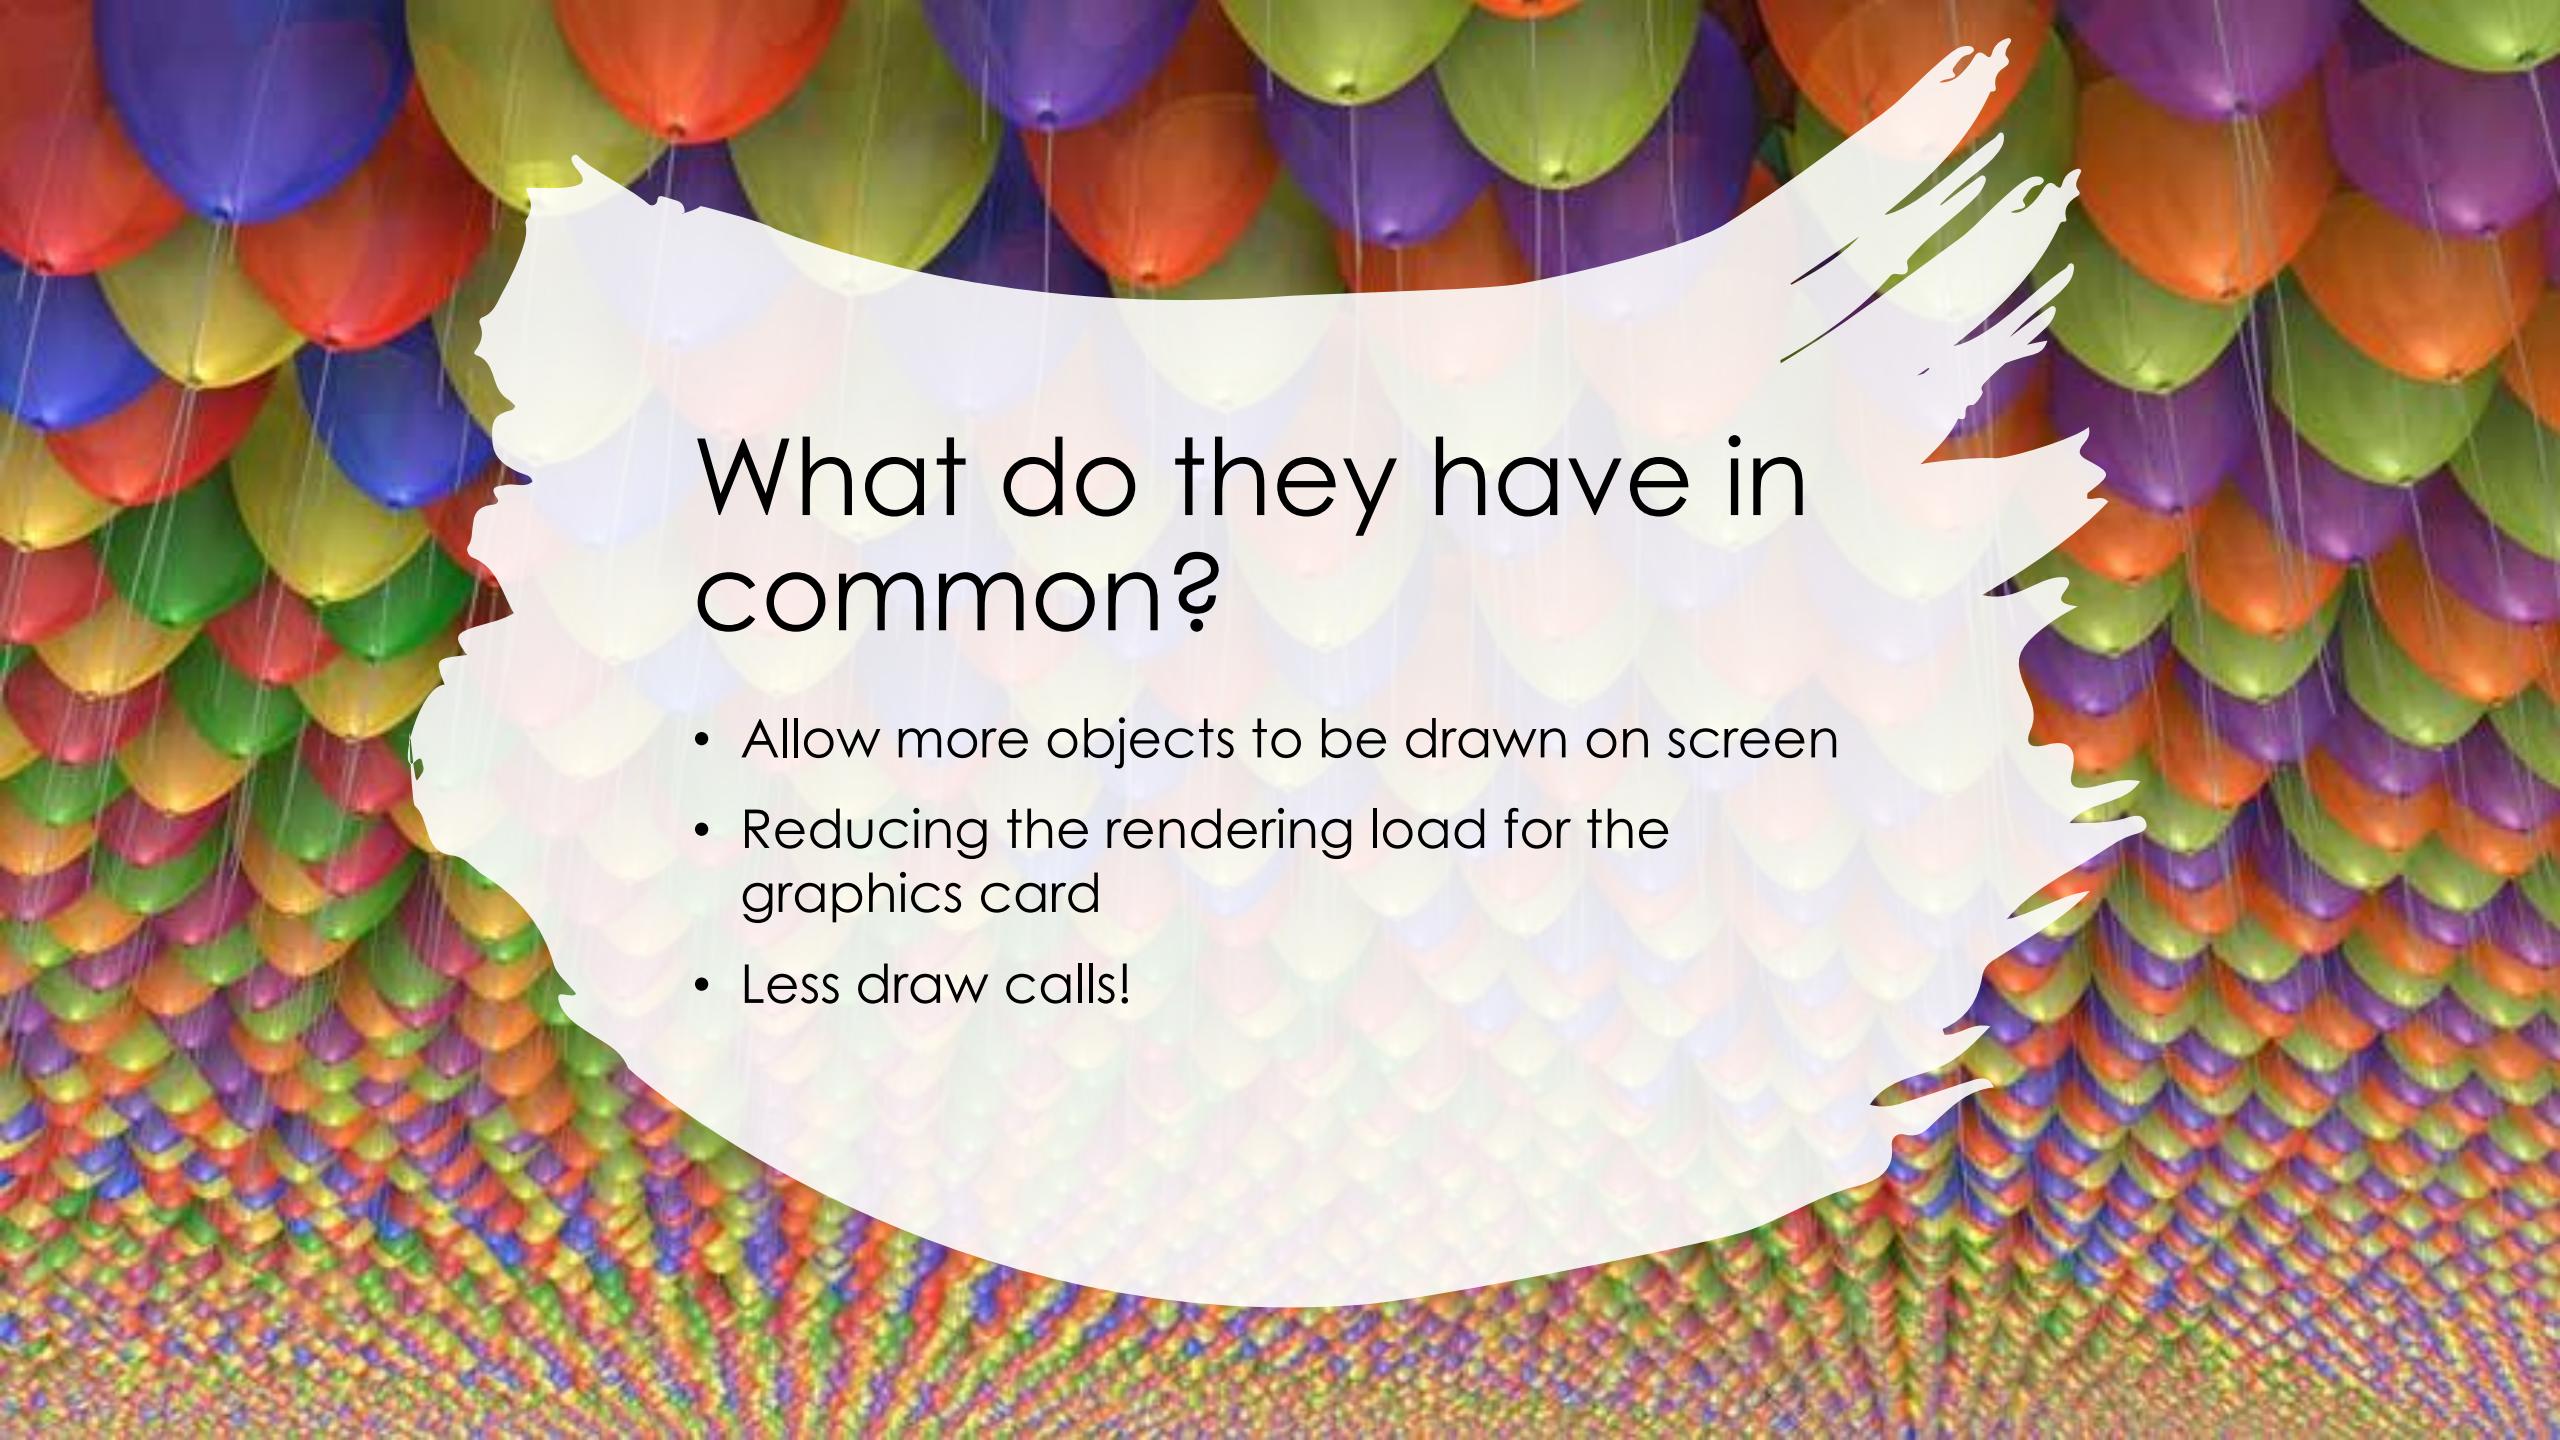
- A Sprite Atlas is an Asset that consolidates several Textures into a single combined Texture
- Unity can call this single Texture to issue a single draw call instead of multiple draw calls to access the packed Textures all at once at a smaller performance overhead
- In addition, the Sprite Atlas API provides you with control over how to load the Sprite Atlases at your Project's run time





Batching vs Instancing

WHAT IS THE DIFFERENCE?



What do they have in common?

- Allow more objects to be drawn on screen
- Reducing the rendering load for the graphics card
- Less draw calls!

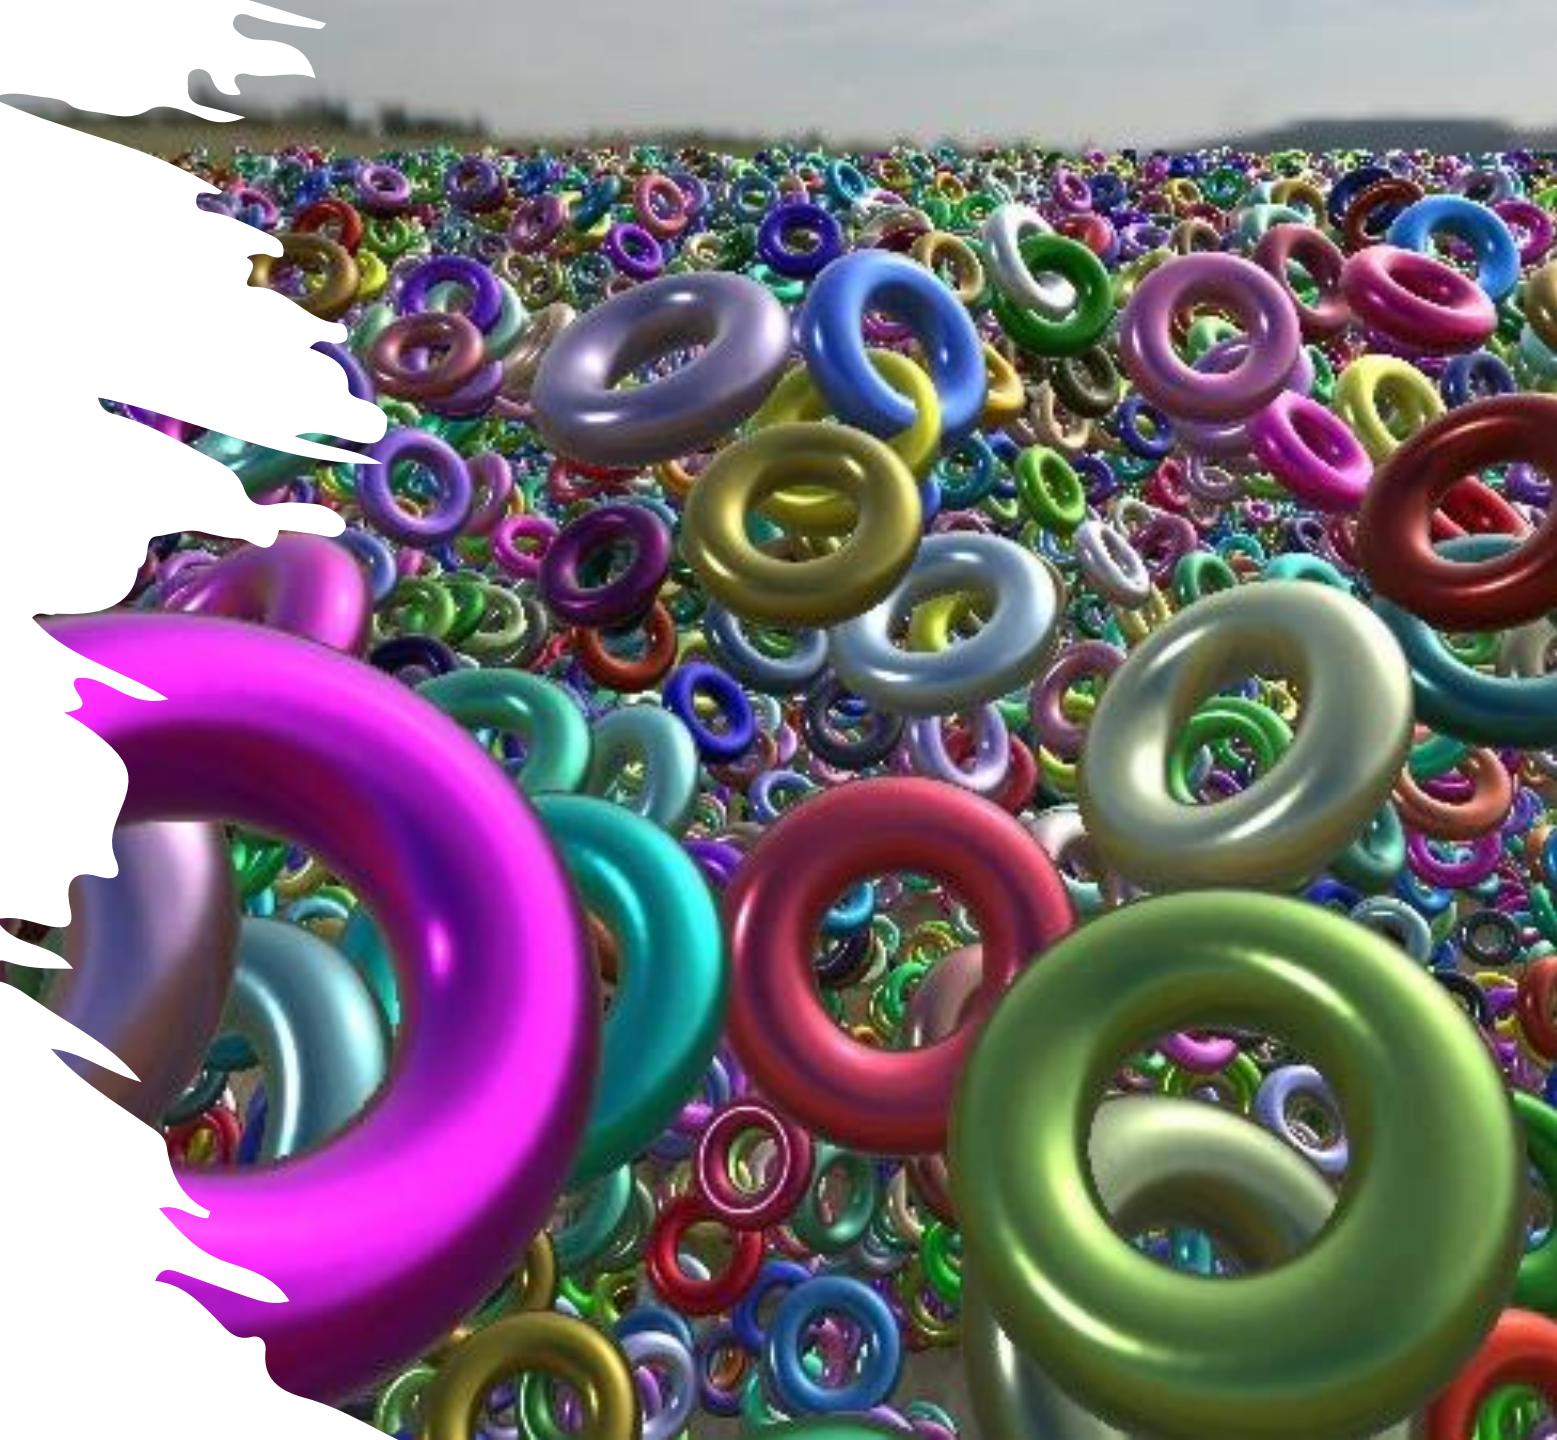
Draw Call Batching in Unity

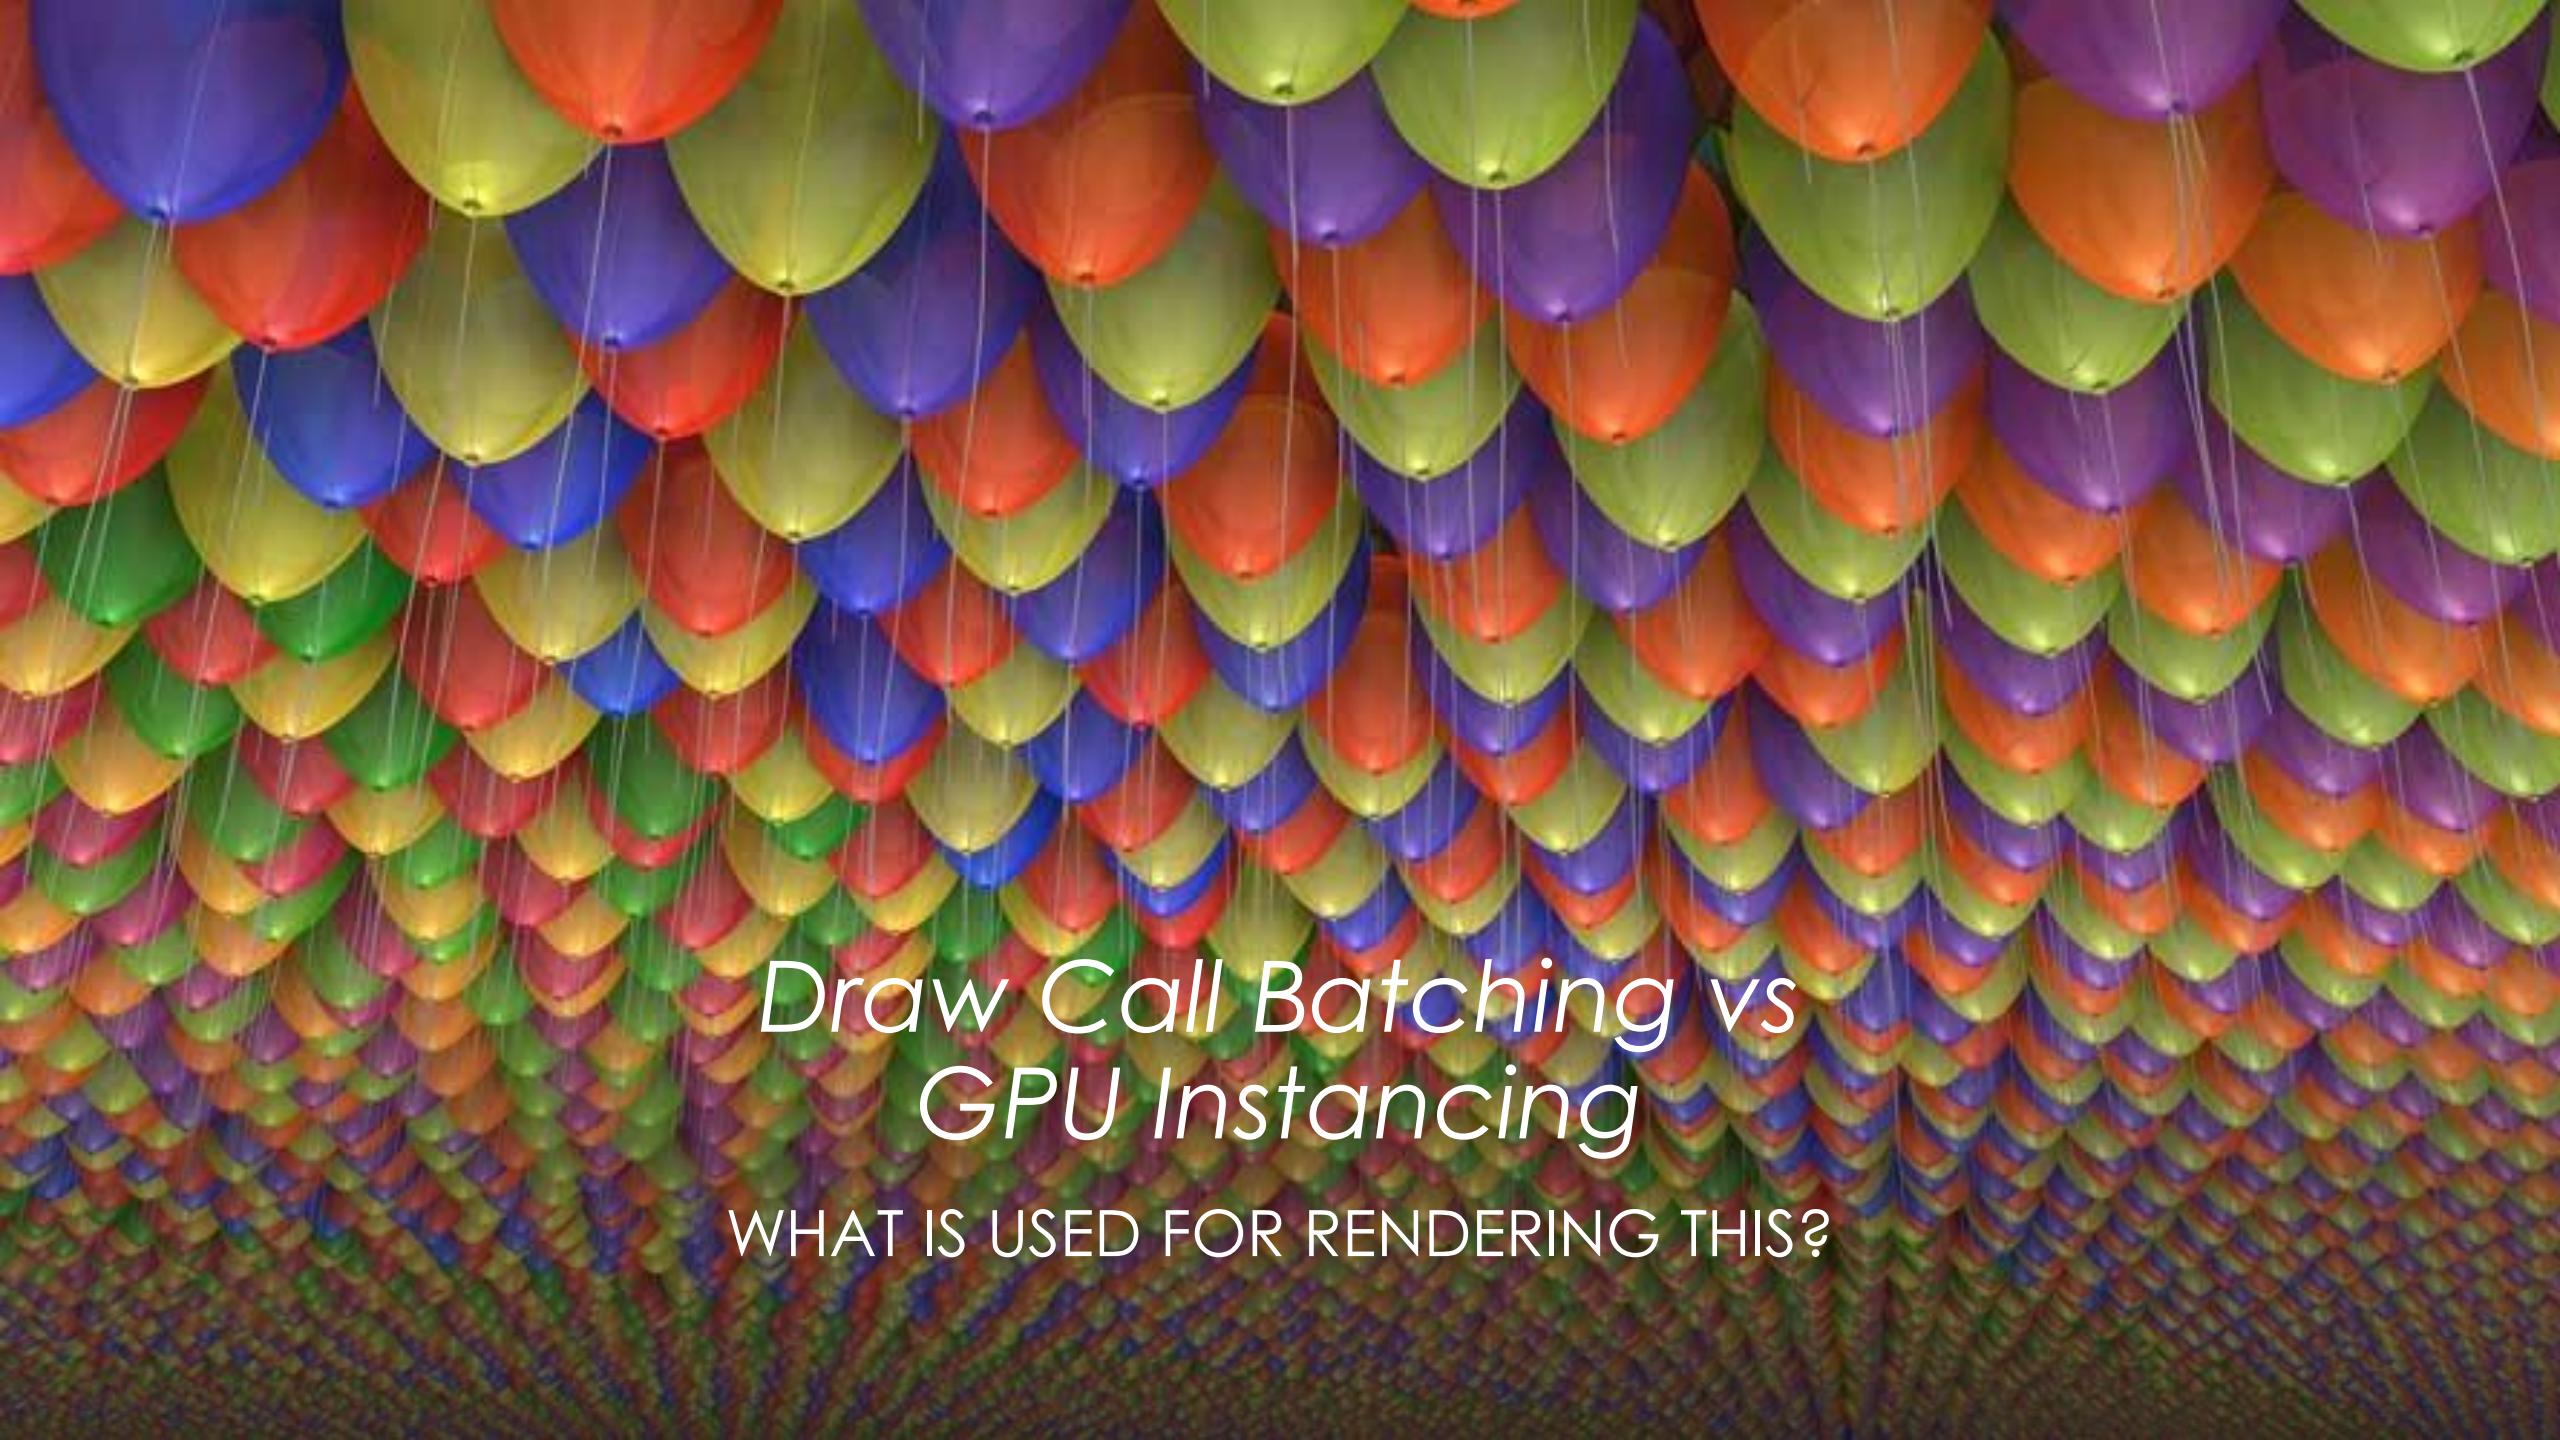
- Batching means to store multiple assets into one GPU buffer
- All geometry properties are merged into buffers and sent to the graphics card
- A buffer can max contain 65.536 vertices
- All assets in a batch must share same material, but using atlas texture can add variety
- Skinned mesh can't be batched, due to CPU transformation processing
- Unity only batches objects of the same type



Draw Call GPU Instancing

- Instancing means to instance a GPU buffer multiple times in one draw call
- An instance can be transformed by offsets to variety to location, rotation and scale
- An instance can max contain 65.536 vertices
- All assets must share same materials but using atlas texture can add variety
- Skinned mesh is not a problem





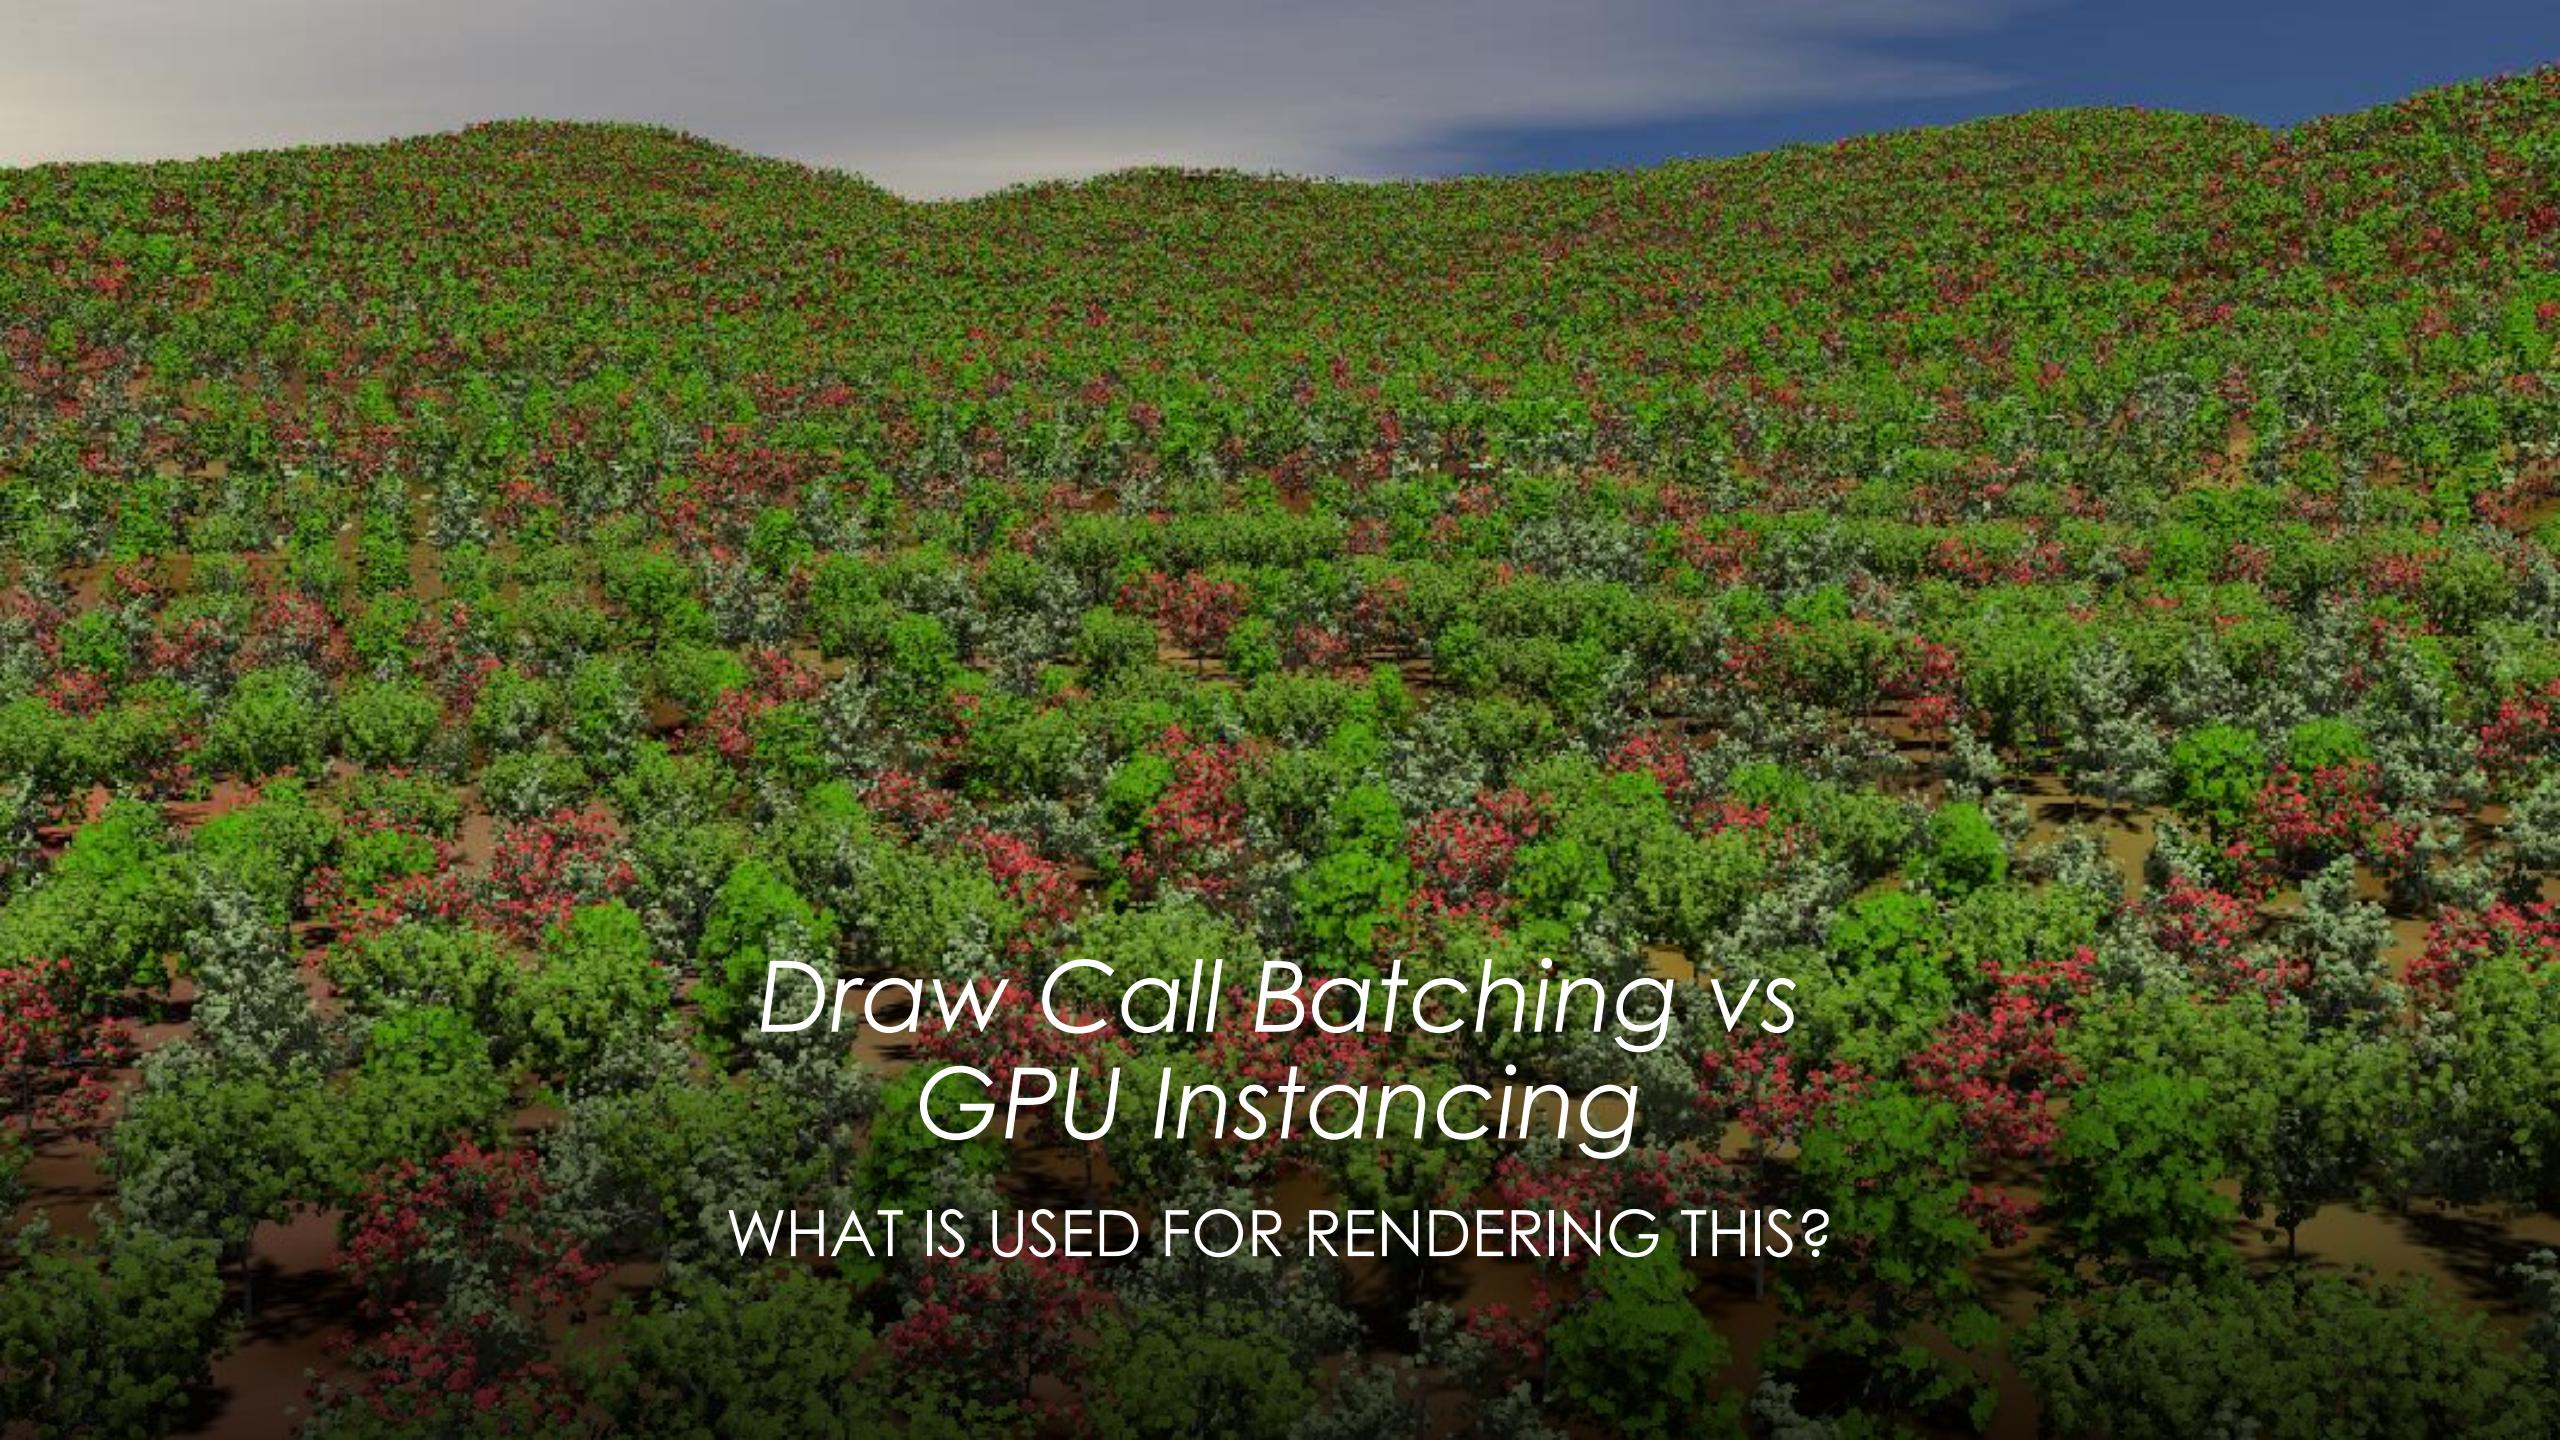
*Draw Call Batching vs
GPU Instancing*

WHAT IS USED FOR RENDERING THIS?



Draw Call Batching vs GPU Instancing

WHAT IS USED FOR RENDERING THIS?

The background of the slide features a dense, sprawling forest of trees with green and reddish-brown foliage, set against a dark, overcast sky.

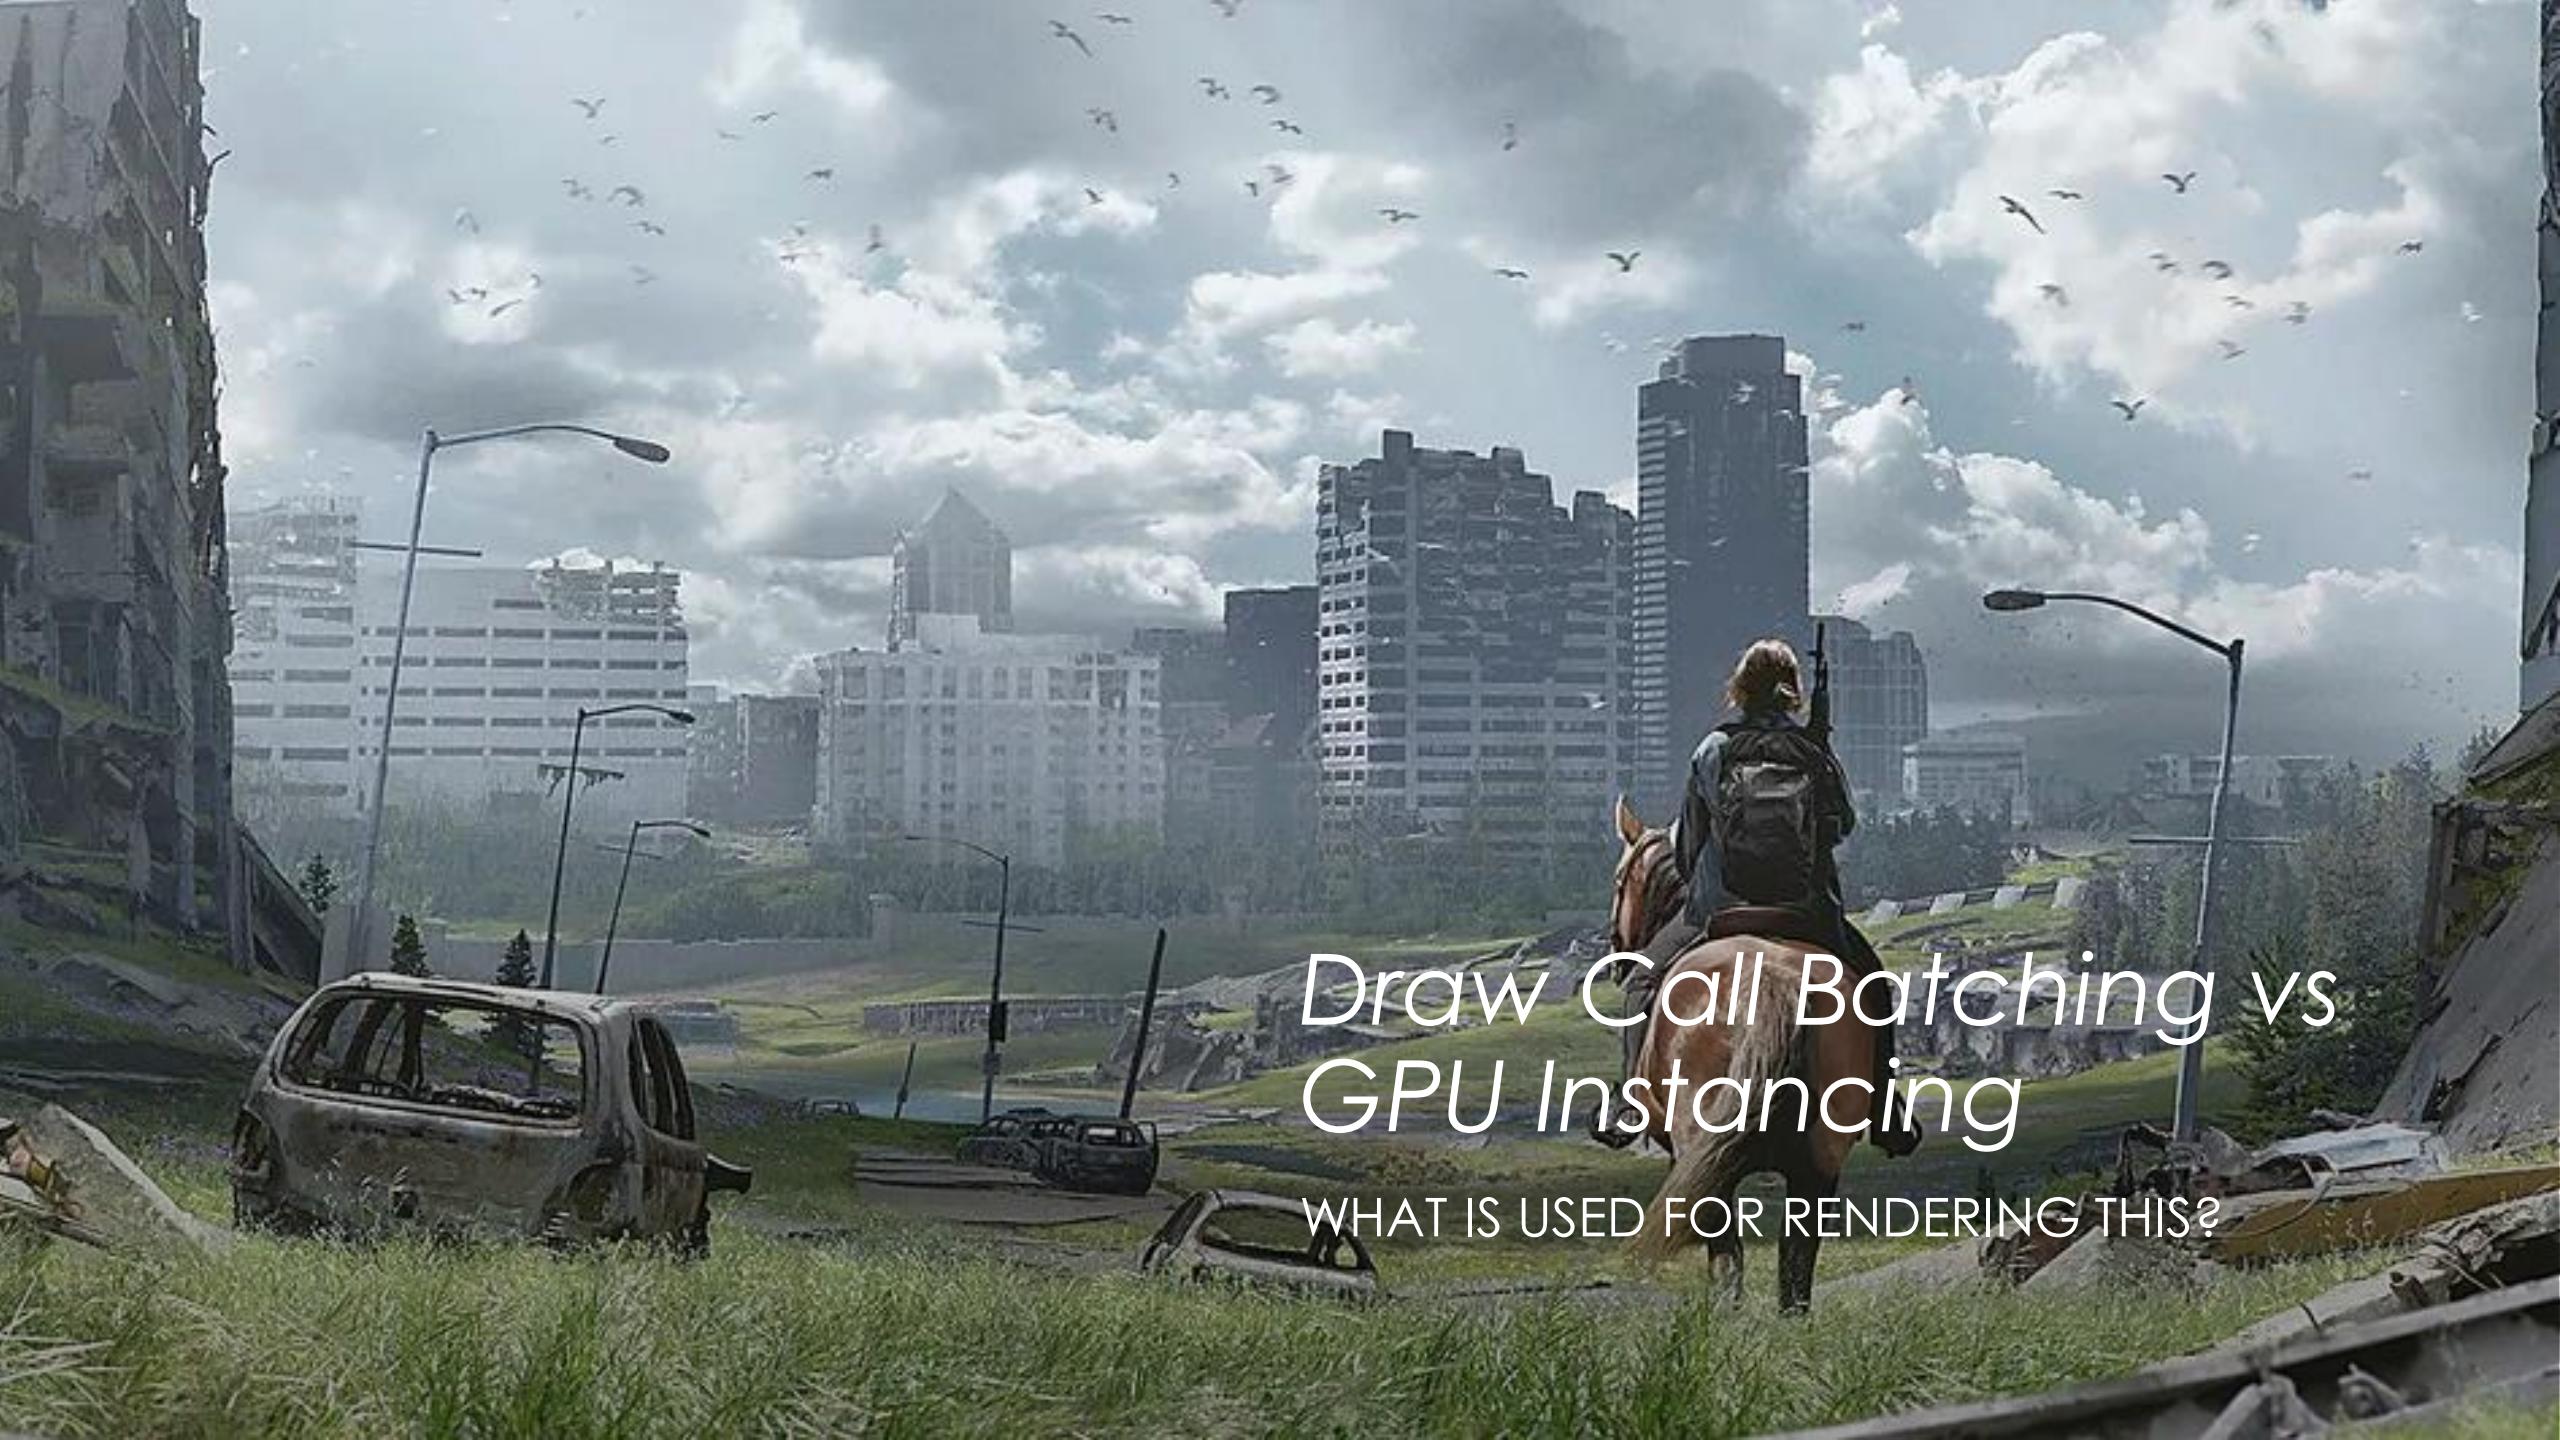
Draw Call Batching vs GPU Instancing

WHAT IS USED FOR RENDERING THIS?



Draw Call Batching vs GPU Instancing

WHAT IS USED FOR RENDERING THIS?

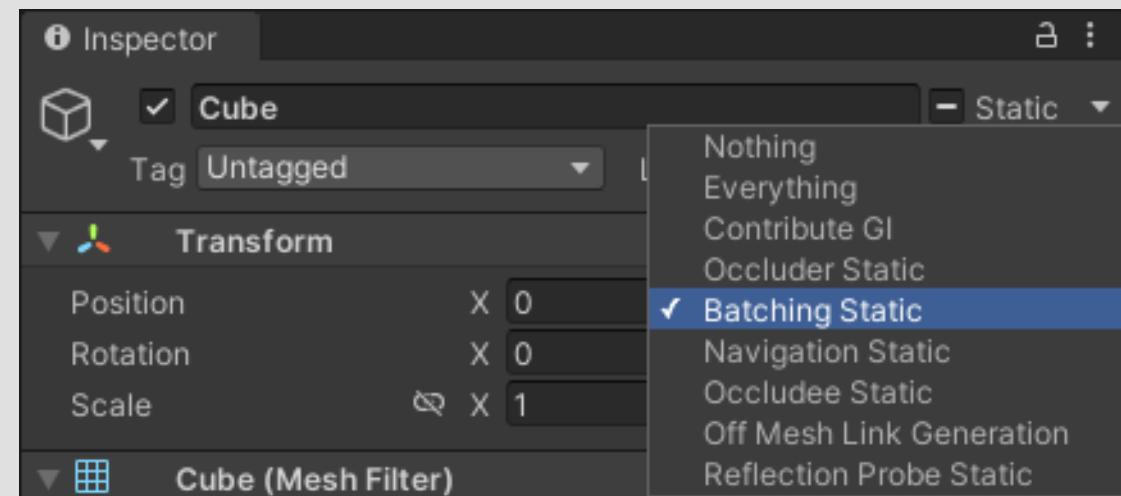
A wide-angle photograph of a desolate, post-apocalyptic urban environment. In the foreground, a woman with long brown hair, wearing a dark jacket and jeans, sits on a brown horse, looking out over the city. She is positioned on the right side of the frame. To her left, a rusted, abandoned car sits on a grassy hillside. The city skyline in the background is filled with numerous skyscrapers of varying heights, all appearing somewhat dilapidated or covered in dust. The sky is filled with a large number of birds flying in various directions. The overall atmosphere is one of a world that has suffered a major catastrophe.

Draw Call Batching vs GPU Instancing

WHAT IS USED FOR RENDERING THIS?

Batching Static Meshes in Unity

- Unity automatically batches the specified static meshes into the same draw call if they fulfill the criteria described in the common usage information



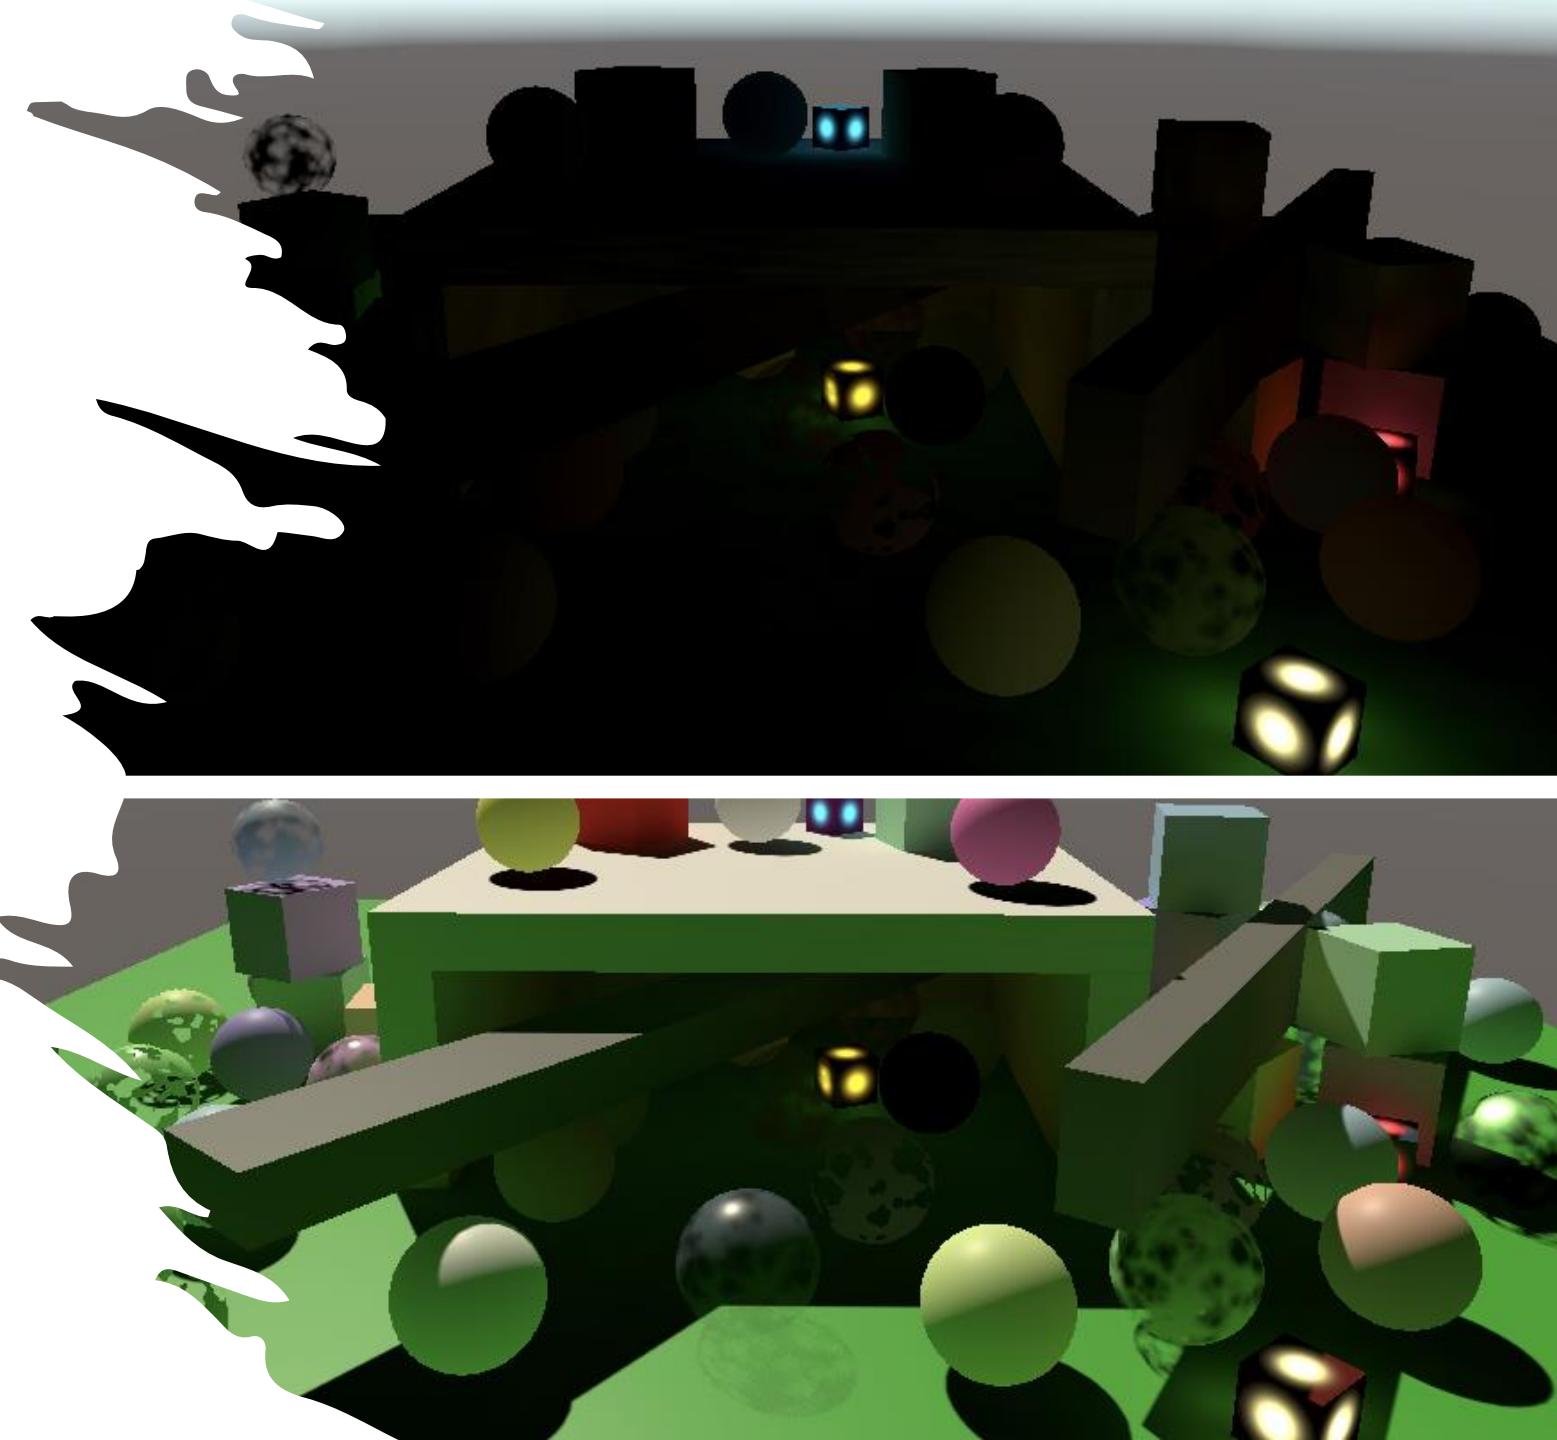
GPU Instancing in Unity

- To use GPU instancing for a material, select the Enable GPU Instancing option in the Inspector
- GPU instancing supports Unity's Baked Global Illumination system
- Unity Standard Shaders and surface shaders support GPU instancing and Unity's Baked Global Illumination system by default



Light Maps & Probes

- Bake lighting when possible to boost performance
- Baked lighting can be mixed with real-time lighting



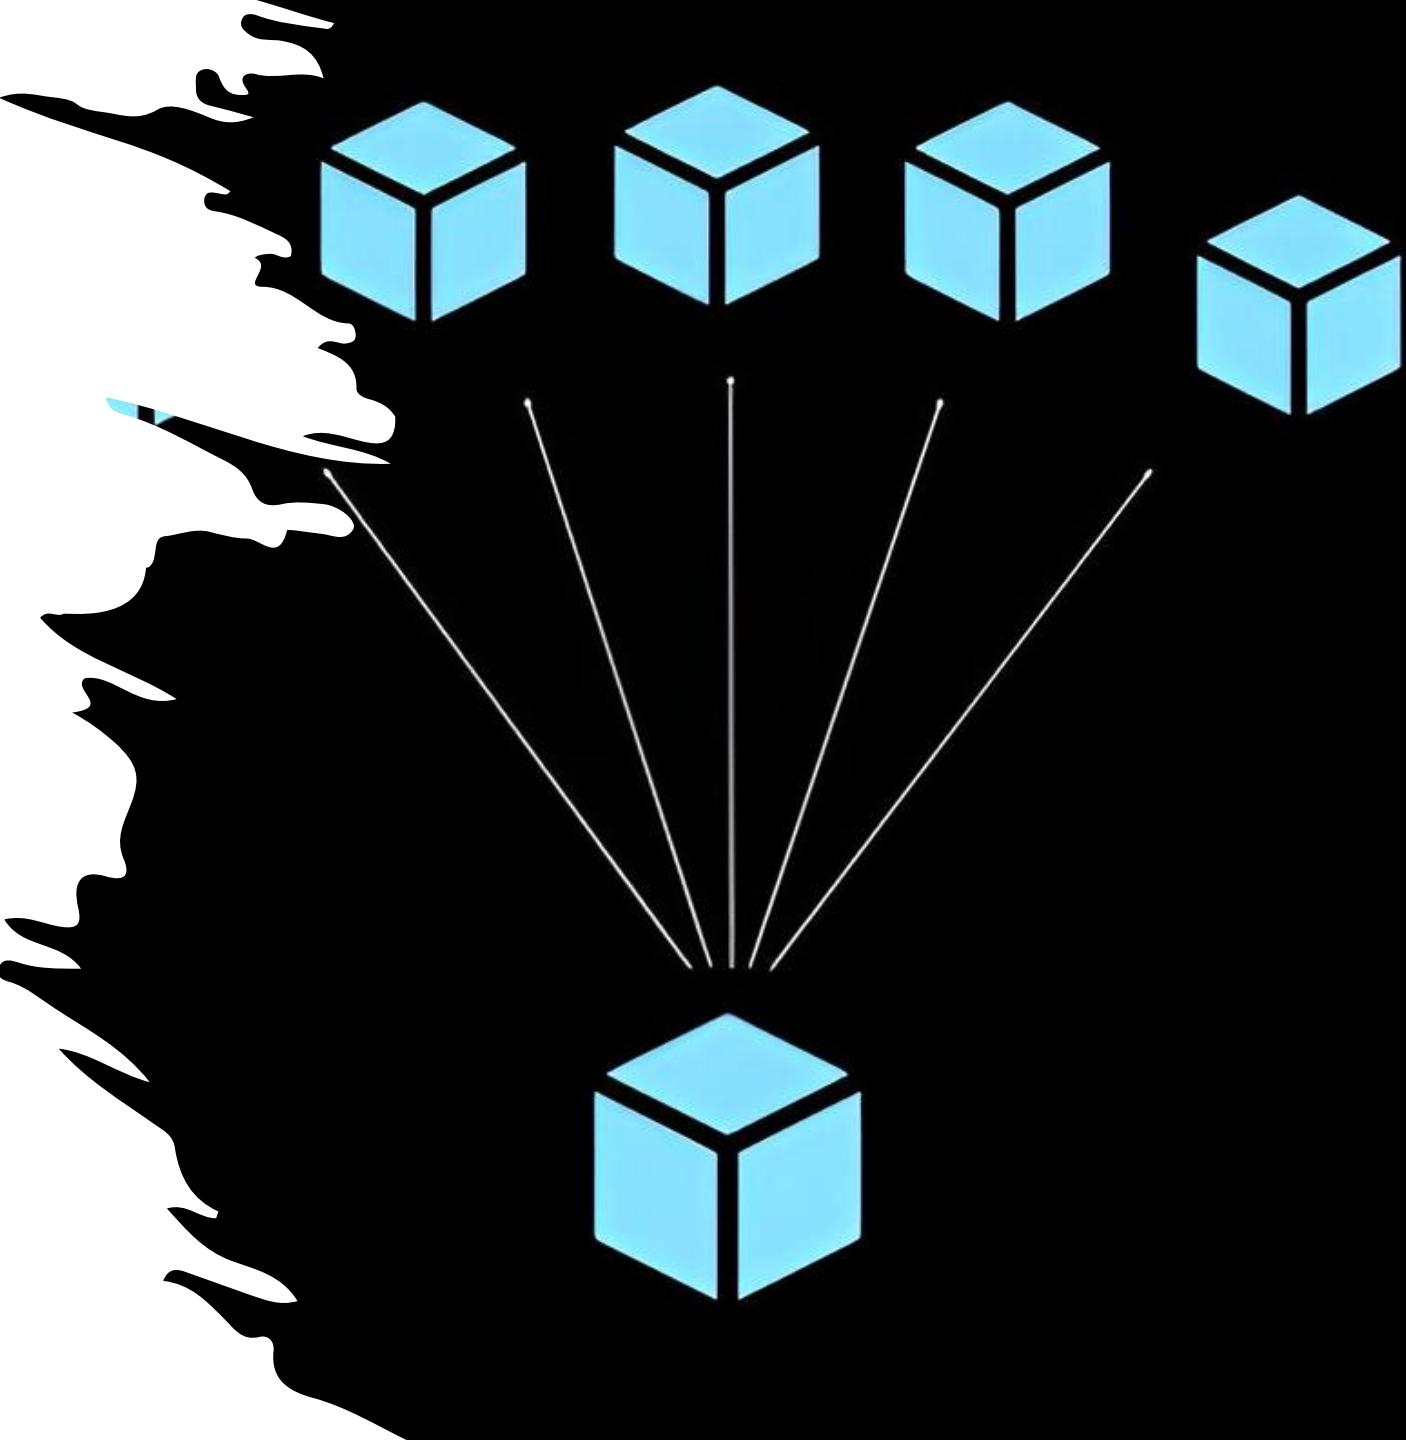
Occlusion Culling

- Only render objects within field of view by splitting up your game world into sections
- Rendering objects that the player cannot see would be a waste of resources



Prefab Workflow

- How many are using prefabs?
- Generally speaking, always use prefab workflow!
- Does Prefab Workflow improve performance?
- Loading time?



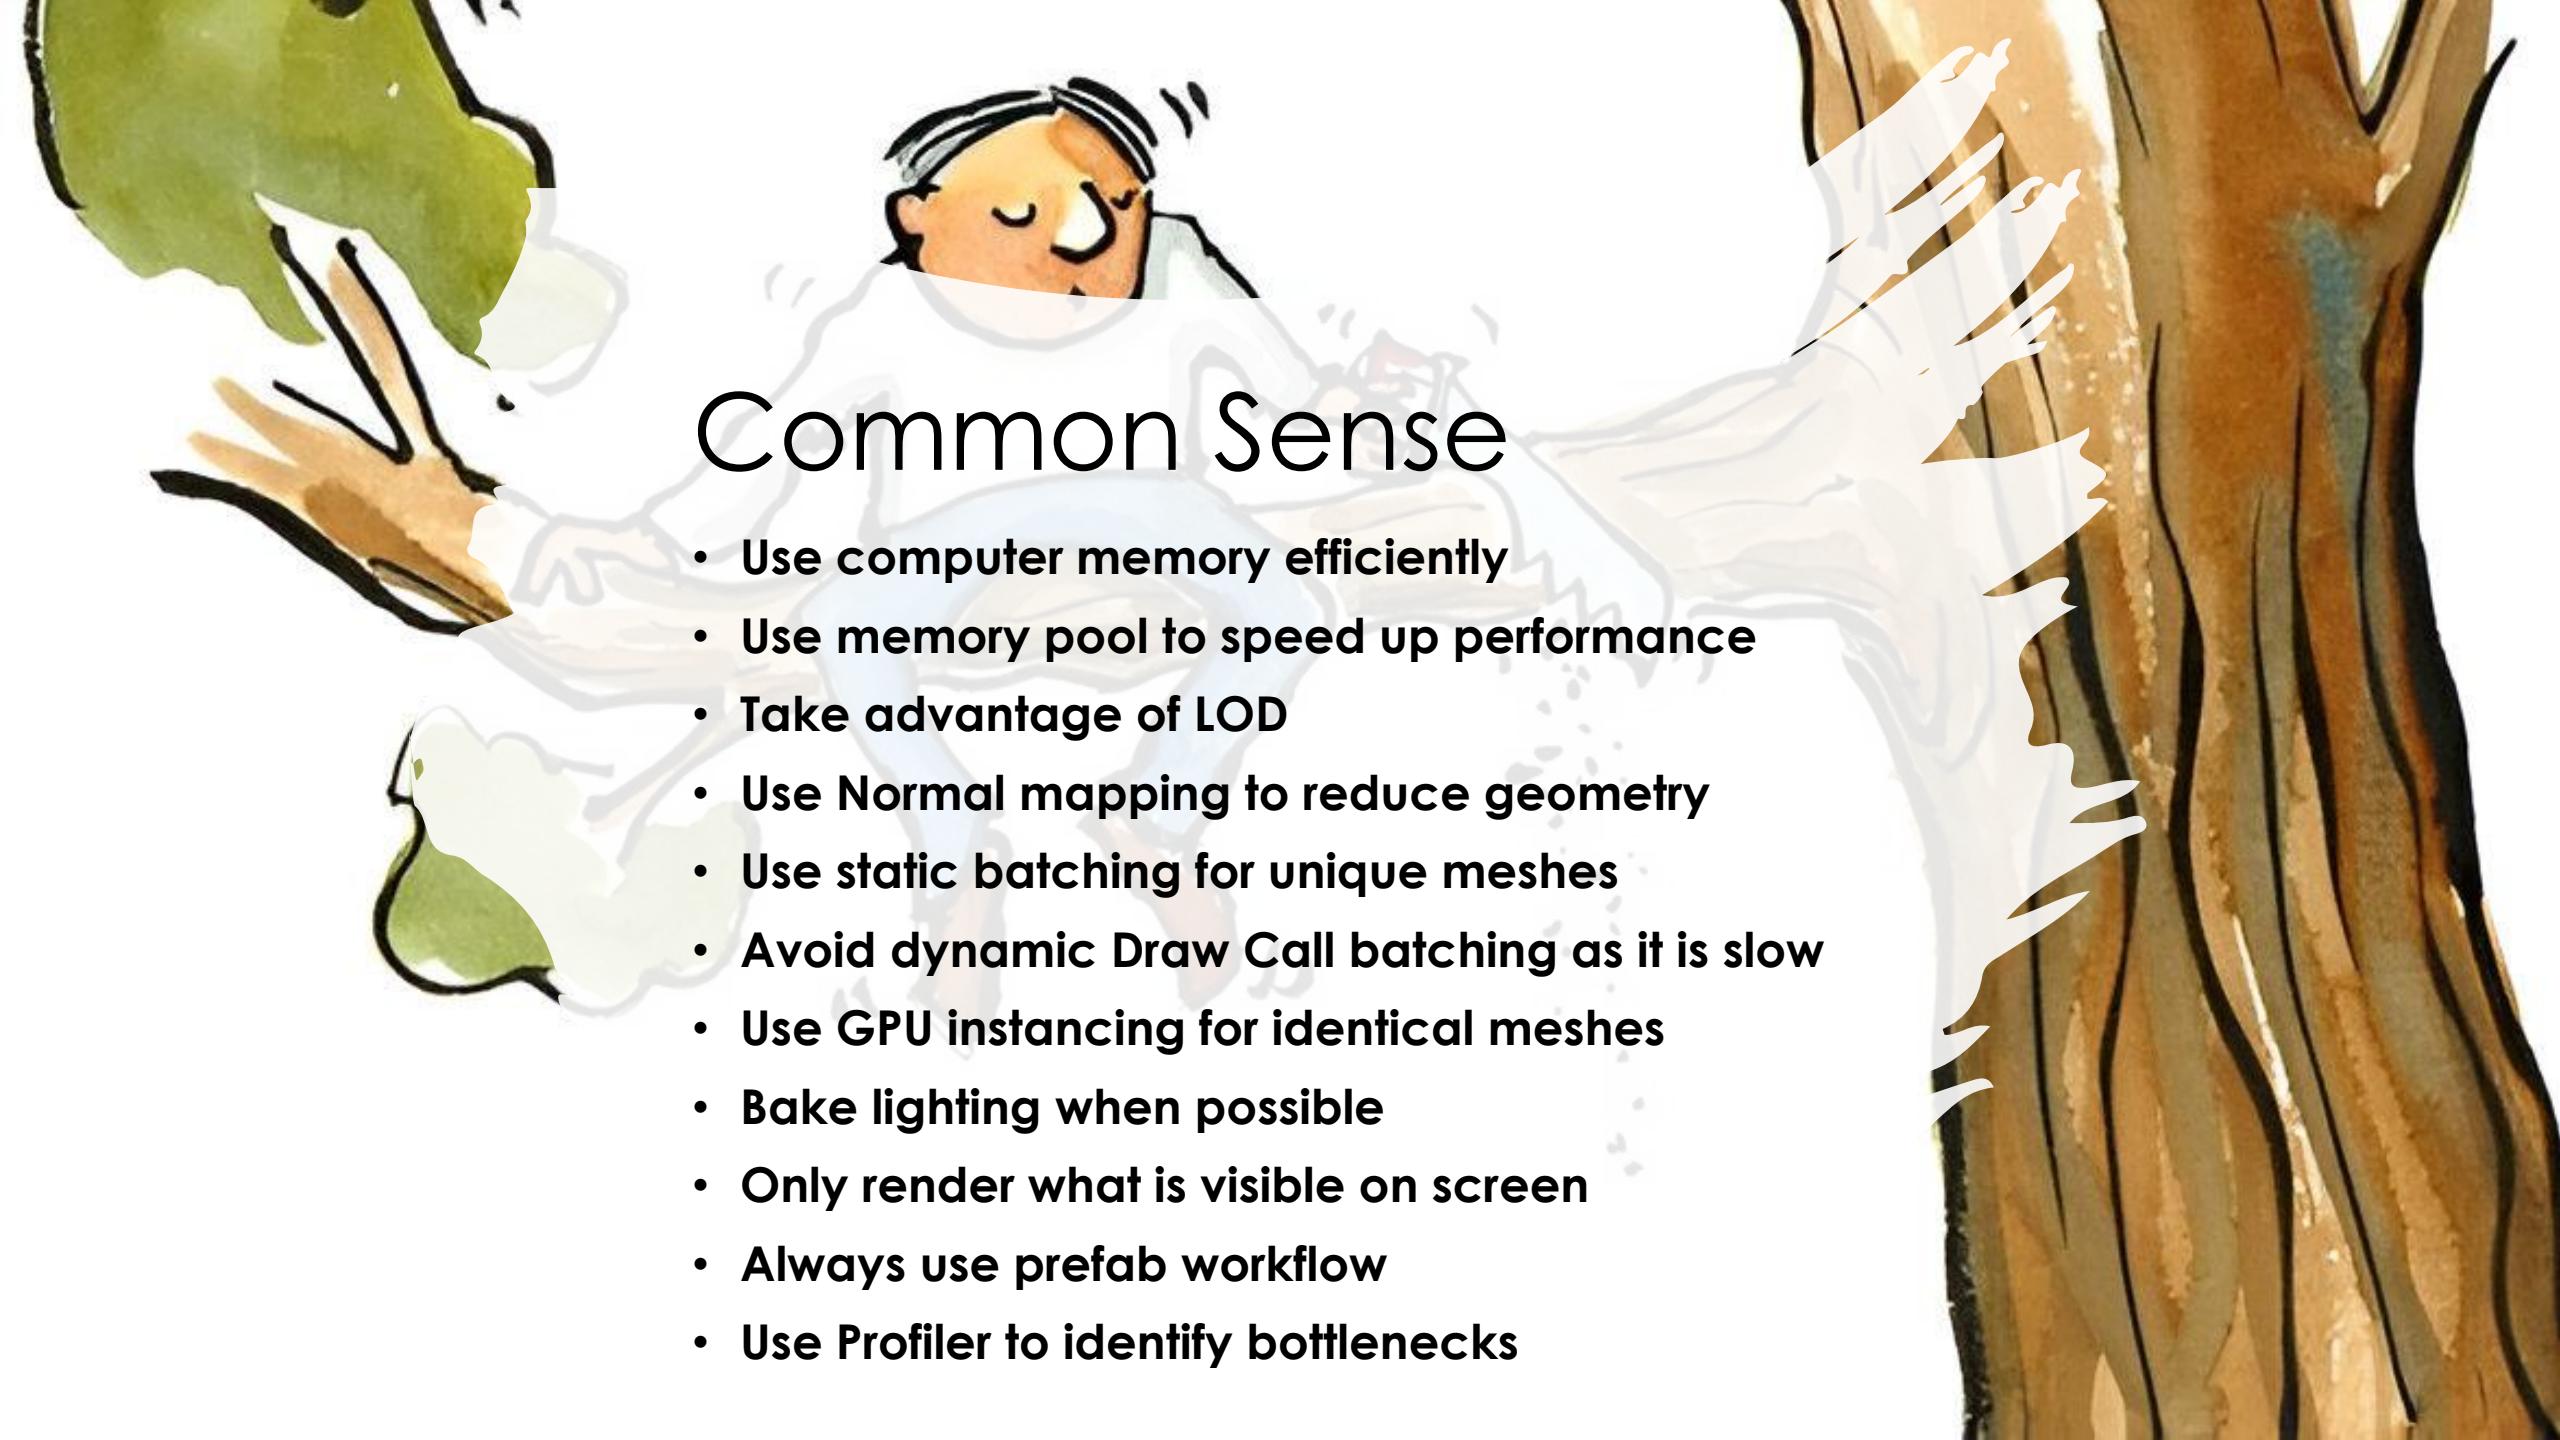


Unity Profiler

USE PROFILER TO IDENTIFY PERFORMANCE BOTTLENECKS

Total	Self ms	CallTime ms	GC ms	Alloc ms	Self ms
56.8%	0.4%	1	0 B	2.06	0.01
39.5%	39.5%	1	0 B	1.43	1.43
1.8%	1.8%	1	0 B	0.06	0.06
0.5%	0.0%	1	0 B	0.02	0.00
0.2%	0.2%	1	0 B	0.01	0.01
0.1%	0.0%	1	68 B	0.00	0.00
0.1%	0.0%	2	0 B	0.00	0.00
0.1%	0.1%	1	0 B	0.00	0.00
0.0%	0.0%	1	0 B	0.00	0.00

Select Line for per-object breakdown



Common Sense

- Use computer memory efficiently
- Use memory pool to speed up performance
- Take advantage of LOD
- Use Normal mapping to reduce geometry
- Use static batching for unique meshes
- Avoid dynamic Draw Call batching as it is slow
- Use GPU instancing for identical meshes
- Bake lighting when possible
- Only render what is visible on screen
- Always use prefab workflow
- Use Profiler to identify bottlenecks